



Building Microservices with Containers

Introduction

Agenda

- Understanding Microservices
- Using Git in DevOps
- Managing Docker Overview
- The Role of Containers in Orchestration
- Understanding Kubernetes
- Creating Container Based Microservices

Poll Question 1

Which of the following topics are most interesting for you? (Choose 3)

- Understanding Microservices
- Using Git in DevOps
- Managing Docker Overview
- The Role of Containers in Orchestration
- Understanding Kubernetes
- Creating Container Based Microservices

Poll Question 2

Rate your knowledge/experience about containers

- none
- minimal
- just attended a basic course
- some working experience
- good working experience
- lots of working experience

Poll Question 3

Rate your knowledge/experience about kubernetes

- none
- minimal
- just attended a basic course
- some working experience
- good working experience
- lots of working experience

Poll Question 4

Which job title applies to you best?

- sysadmin
- devops
- developer
- security engineer
- DBA
- network engineer
- management
- other (use group chat to specify)

Poll Question 5

Which part of the world are you from?

- Europe
- Netherlands
- Africa
- North/Central America
- South America
- India
- Asia
- Australia/Pacific

Course Setup

- A running environment with Kubernetes (use cloud hosted for easy access, or minikube for local access)
- Course github repository is available at <https://github.com/sandervanvugt/microservices>

A large, light gray play button icon is positioned on the left side of the slide. It consists of a white right-pointing triangle centered within a series of concentric gray circles.

Building Microservices with Containers

1. Understanding Microservices

What are Microservices?

- Microservices are grains of application code that run minimal functionality that can be isolated from other grains of application code
- The different grains are loosely coupled to each other
- The different grains are independently developed and maintained

Microservices benefits

- When broken down in pieces, applications are easier to build and maintain
- Each piece is independently developed, and the application itself is just the sum of all pieces
- Smaller pieces of application code are easier to understand
- Developers can work on applications independently
- If one component fails, you spin up another while the rest of the application continues to function
- Smaller components are easier to scale
- Individual components are easier to fit into continuous delivery pipelines and complex deployment scenarios

Understanding Containers and Microservices

- A container is an application that runs based on a container image. The container image is a lightweight standalone executable package of software that included everything that is needed to run a application
- Because containers are lightweight and include all dependencies required to run an application, containers have become the standard for developing applications
- As containers are focusing on their specific functionality, they are perfect building blocks for building microservices
- The main requirement is that the different components are connected to eachother the right way, by providing static values instead of variables
- These static values can be provided in a very flexible way by using container orchestration platforms such as Kubernetes

How to Break up Monolithic Applications

- While breaking up monolithic applications, several changes are often applied
 - Connection parameters to database and middleware need to be changed from hard-coded to variables that can be managed in a flexible way
 - In web applications, application calls need to be changed to fixed public DNS host names
 - Security needs to be modified, to ensure that one application is allowed to access other applications. This is known as cross-origin resource sharing
- Breaking up Monolithic applications can be made easier using different solutions that are provided in Kubernetes and other orchestration platforms

What is RESTful API

- REST stands for REpresentational State Transfer and it's an architectural style for distributed systems
- REST is based on six guiding principles
 - Client-server: the client interface is separated from the server part to increase portability
 - Stateless: the server keeps no information about the state of the request, and each request must contain all the required information
 - Cacheable: the server must indicate if data in the response can be cached
 - Uniform interface: the interface to the server is well defined and uniform
 - Layered system: an architecture should be designed as multiple hierarchical layers, where a component can not see beyond its own layer
 - Code on demand: REST allows client functionality to be extended by downloading and executing code as applets or scripts

What is RESTful API?

- The RESTful API defines resource methods to perform a desired transition to interact with web services
- Commonly, the HTTP GET/PUT/POST/DELETE methods are used for this purpose
- Key feature is that the RESTful API should be entered without prior knowledge beyond the initial URI
- RESTful API is the standard in microservices, as it provides an easy to use, stateless and uniform way for different parts of the microservice to interact

What is CI/CD

- CI/CD stands for continuous integration / continuous delivery/deployment
- CI/CD is a method to deliver apps in a flexible and automated way
- CI/CD makes integration of new code in a solution easy
- CI/CD introduces ongoing automation and monitoring throughout the lifecycle of applications, from integration to delivery and deployment

The Role of CI/CD in Microservices

- In a microservices oriented way of working, it is important to implement changes in an easy and non-disruptive way
- CI/CD is used to guarantee this can happen. The code typically originates in a Git repository, from which it can be deployed by using automation solutions like Dockerfile or OpenShift, which on its turn are hosted in a Git repository as well
- By hosting everything in Git repositories, it's easy to manage development and version differences

What is DevOps?

- In DevOps Developers and Operators work together on implementing new software and updates to software
- The purpose of DevOps is to reduce the time between committing a change to a system and the change being placed in production
- DevOps is microservices oriented by nature, as more smaller projects are so much easier to manage than one monolithic project
- In DevOps, CI/CD pipelines are commonly implemented, using anything from simple GitHub repositories, up to advanced CI/CD oriented software solutions such as Jenkins and OpenShift

Understanding the goals of this course

- In this course you'll learn how to build Microservices using containers
- We'll explore the devops cycle, starting with the Git repositories and the Docker files, continuing with a strong focus on containers and their workings, and ending in the use of Microservices in the Kubernetes orchestration platform
- This course does NOT focus on the programmatic parts of microservices, so knowledge of specific programming languages is not required
- The main goal in this course is to teach students how to build a microservices based infrastructure based on orchestrated containers

A large, light gray play button icon with a white triangle pointing right, centered within a circle. The circle has a thick white border and is set against a dark gray background.

Building Microservices with Containers

2. Using Git in DevOps

Using Git in a Microservices Environment

- Git can be used for version control and cooperation between different developers and teams
- Using Git makes it easy to manage many changes in an effective way
- In a Microservices way of working, this is exactly what is needed
- For that reason, Git and Microservices are a perfect match

Understanding Git

- Git is a version control system that makes collaboration easy and effective
- Git works with a repository, which can contain different development branches
- Developers and users can easily upload as well as download new files to and from the Git repository
- To do so, a Git client is needed
- Git clients are available for all operating systems
- Git servers are available online, and can be installed locally as well
- Common online services include GitHub and GitLabs

Git Client and Repository

- The Git repository is where files are uploaded to, and shared to other users
- Individual developers have a local copy of the Git repository on their computer and use the **git** client to upload and download to and from the repository
- The organization of the Git client lives in the `.git` directory, which contains several files to maintain the status

Understanding Git Workflow

- To offer the best possible workflow control, A Git repository consists of three trees maintained in the Git managed directory
 - The *working directory* holds the actual files
 - The *Index* acts as a staging area
 - The *HEAD* points to the last commit that was made
- The workflow starts by creating new files in the working directory
- When working with Git, the **git add** command is used to add files to the index
- To commit these files to the head, use **git commit -m "commit message"**
- Use **git add origin https://server/reponame** to connect to the remote repository
- To complete the sequence, use **git push origin master**. Replace "master" with the actual branch you want to push changes to

Creating a Git Repository

- Create the repository on your Git server
- Set your user information
 - **git config --global user.name "Your Name"**
 - **git config --global user.email "you@example.com"**
- Create a local directory that contains a README.md file. This should contain information about the current repository
- Use **git init** to generate the Git repository metadata
- Use **git add <filenames>** to add files to the staging area
- From there, use **git commit -m "commit message"** to commit the files. This will commit the files to HEAD, but not to the remote repository yet
- Use **git remote add origin https://server/reponame**
- Push local files to remote repository: **git push -u origin master**

Using Git Repositories

- Use **git clone https://gitserver/reponame** to clone the contents of a remote repository to your computer
- To update the local repository to the latest commit, use **git pull**
- Use **git push** to send local changes back to the git server (after using **git add** and **git commit** obviously)

Uploading Changed Files

- Modified files need to go through the staging process
- After changing files, use **git status** to see which files have changed
- Next, use **git add** to add these files to the staging area; use **git rm <filename>** to remove files
- Then, commit changes using **git commit -m "minor changes"**
- Synchronize, using **git push origin master**
- From any client, use **git pull** to update the current Git clone

A large, light gray play button icon is positioned on the left side of the slide. It consists of a white right-pointing triangle centered within a series of concentric gray circles.

Building Microservices with Containers

3. A Micro-introduction to Docker

Containers Defined

- A container is an application with all of its dependencies included
- Containers always start an application, they don't sit down and wait for you to tell them what to do
- To run a container, a container runtime is needed. This is the layer between the host operating system and the container itself
- Different solutions exist to run containers
 - Docker: owns more than 80% of the market and is the de facto standard
 - LXC: Linux native containers
 - CoreOS rkt AKA podman: works with Pods and natively integrates with Kubernetes

Understanding Container Types

- System containers behave like virtual machines. They go through a complete boot process and provide VM-typical services like SSH and logging
- Application containers are used to start just one application. Application containers have become the default
- Orchestration solutions like Kubernetes focus on managing application containers and may become confused while working with system images as they don't have a default application to start

Understanding a Perfect Match

- Microservices focus on developing minimal pieces of code and joining them
- Containers are focussing on developing minimal running application components
- So both have the same objective in mind
- The only thing that needs to be added, is a layer that connect containers together
- This is going to be done by the container orchestration layer

Containers and Images

- A container is a running instance of an image
- The image contains the application code, language runtime and libraries
- External libraries such as libc are typically provided by the host operating system, but in a container is included in the images
- The container image is a read-only instance of the application that just needs to be started
- While starting a container, it adds a writable layer on the top to store any changes that are made while working with the container

A large, light gray play button icon is positioned on the left side of the slide. It consists of a white right-pointing triangle centered within a series of concentric gray circles.

Building Microservices with Containers

4. The Role of Container Orchestration

Understanding Enterprise Container Requirements

- Containers are based on the Container Runtime, a solution that has nothing to connect multiple containers together in a Microservice architecture
- Containers need to be distributed, scheduled and load balanced
- Apart from that, High Availability is required
- As well as an easy solution to provide updates without downtime
- To run containers as microservices, additional platform components are required as well, such as software defined networking and software defined storage

Understanding Microservices Platform Requirements

- To understand Microservices platform requirements, you need to understand the typical Microservices application
- Microservices applications typically drill down to a database part and an accessibility part, using REST or Web to provide accessibility

From Microservice to Platform

- Based on the Microservice overview on the previous slide, the following minimal platform requirements can be defined
 - Different databases must be running as connected applications
 - Front-end services need to be added to that, and add an accessibility layer as well
 - At the user side, accessibility must be added to different types of user requests
- Add some scalability to this, as well as some high availability, and you'll have the basic platform requirements
- And optionally, add integration of a CI/CD pipeline as well to make the cycle from source code to application complete

Understanding the Orchestration Landscape

- Kubernetes is the leading solution and the open source upstream for many other platforms
- Docker Swarm was developed by Docker Inc and offered as Docker Enterprise - now a part of Mirantis
- Red Hat OpenShift is based on the OpenShift Kubernetes Distribution (OKD) and offers Kubernetes with strong developed CI/CD

Understanding the Leading Position of Kubernetes

- Kubernetes is the leading technology that is in nearly all container orchestration platforms
- This is because of its origins, coming from Google Borg
- Google donated the Kubernetes specifications to the open source community after running Borg internally for over a decade
- As a result, a free, stable and open source project was introduced, where all vendors could base their own solution upon
- This rocked the world of container orchestration, and resulted in Kubernetes being the leading platform for orchestration

Understanding Kubernetes Delivery Options

- Kubernetes as a managed service in public and private cloud
 - Amazon EKS
 - Google Cloud Platform GKE
 - Azure AKS
 - OpenStack Magnum
- Kubernetes as an on-premise installation using a Kubernetes distribution
- Kubernetes as a test-drive platform, mainly developed to learn Kubernetes

Understanding Kubernetes Distributions

- When using on premise Kubernetes, something needs to be installed
- Open Source Kubernetes can easily be installed using **kubeadm** on any supported platform
- Kubernetes Distributions offer the additional value of providing support and on occasion more stable code
- Most on-premise Kubernetes distributions can be used in cloud as well

Common Container Distributions

- Canonical Kubernetes: offered by the makers of Ubuntu, relatively pure and based on open source, runs on premise and in cloud
- Rancher: focus on offering Kubernetes as multi-cluster deployments. Available as on-premise and in cloud
- SUSE CaaS platform: offered by SUSE as an on-premise and private cloud solution - support for public cloud is pending
- OpenShift: developed by Red Hat, available on-premise and in cloud. Adds CI/CD features to core Kubernetes functionality

A large, light gray play button icon is positioned on the left side of the slide. It consists of a white right-pointing triangle centered within a series of concentric circles, all rendered in a light gray color.

Building Microservices with Containers

5. Understanding Kubernetes

Getting Started with Minikube

- Minikube is an all-in-one Kubernetes virtual machine
- Using Minikube is recommended as it gives access to all Kubernetes features, without any limitations
- In this course you'll learn how to set up Minikube on a Fedora virtual or physical machine
- For instructions on setting up Minikube on other platforms, got to: <https://kubernetes.io/docs/tasks/tools/install-minikube/>
- To do the hands on parts in the next lessons, it's fine to use any other Kubernetes platform as well

Installing Minikube in Fedora

- Install the most recent version of Fedora Workstation, using 40GB hard disk space, 8GB of RAM (more is recommended), and 2 CPU cores with embedded virtualization enabled
- Notice that Oracle Virtual Box does NOT support embedded virtualization for AMD processors
- Use **git clone** <https://github.com/sandervanvugt/microservices> to access the GitHub repository for this course
- Run the **kube-setup.sh** script from the GitHub repository and follow the instructions it provides

Running Containers in Kubernetes

- Do NOT run naked pods unless if its for testing
- **kubectl run** will run a pod
- **kubectl create deployment** will create a deployment

Understanding Decoupling in Microservices

- In microservices, it's all about independent development cycles, allowing developers to focus on their chunk of code
- To make reusing code easy, separation of static code from dynamic values is important
- This approach of decoupling items should be key in your Microservices strategy
- Kubernetes helps, by adding many resource types that focus on decoupling

Decoupling Resource Types in K8s

- ConfigMap: used for storing variables and config files
- Secrets: used for storing variables in a protected way
- PersistentVolumes: used to refer to external storage
- PersistentVolumeClaims: used to point to the storage type you need to use

Running Containers in Pods and Deployments

- To run containers, the *imperative* approach can be used
 - **kubectl create deployment**
 - **kubectl run**
- Alternatively, the *declarative* approach can be used, where specifications are done in a YAML file
- The declarative approach is preferred in a DevOps environment, as versions of YAML manifests can easily be maintained in Git repositories

Understanding YAML

- YAML is a markup language, where parent-child relations between objects are indicated using indentation
- The objects are key: value type objects
- If a key can have multiple values, each value is indicated with a hyphen
- Use spaces for indentation, using tabs is forbidden
- Use the following in ~/.vimrc to make it easier to create YAML
 - **au! BufNewFile,BufReadPost *.{yaml,yml} set filetype=yaml foldmethod=indent**
 - **autocmd FileType yaml setlocal ts=2 sts=2 sw=2 expandtab**

Understanding the API

- The API server provides access to Kubernetes objects
- The API is also the specification of objects and their properties
- **kubectl** is the key utility to write objects to the API
- The Kubernetes API is extensible, which makes it easy to add new objects
- This is continuously happening, which is also why a new version Kubernetes is released every 3 months

Exploring the API

- **kubectl api-resources** will show the different resources, including the API group they come from
- **kubectl api-versions** will show different objects, the API they come from, as well as the current version
- **kubectl explain** allows you to browse through the different objects in the API and explore how they are constructed
 - This is a key skill for discovering how to set up the YAML code where objects are defined

Essential Troubleshooting

- **kubectl get** gives an overview of different object types and should be used as first step in troubleshooting
- **kubectl describe** explores Kubernetes resources as created in the Kubernetes etcd database
- **kubectl logs** shows the STDOUT for applications running in a container and is useful for application troubleshooting
- **kubectl events** gives an overview of recent events

A large, light gray play button icon is positioned on the left side of the slide. It consists of a white right-pointing triangle centered within a series of concentric circles, all rendered in a light gray color.

Building Microservices with Containers

6. Creating Container Based Microservices

Managing Kubernetes Registries

- If no registry hostname is specified, Kubernetes fetches images from the Docker public registry
- To manage images stored in other registries, refer the complete path to the image: **my.registry.example.com/myimage**
- For using images from private registries, authentication may be required
- Perform the authentication at the container layer, for instance by making a `$HOME/.docker/config.json` file

Understanding imagePullPolicy

- By default, images are fetched from registries
- To use local images, use **imagePullPolicy: IfNotPresent** or **imagePullPolicy: Never**
- For working with local images, you'll need to make sure all cluster nodes have local pre-pulled images
- On the Kubernetes nodes, images are managed by the container layer and typically stored in an image database
- If the container layer is Docker, use **docker images ls** for a list of images
- To set up a private registry, use the procedure offered by the container layer

Understanding Pod Storage

- Volumes are a part of the Pod container spec
- For increased flexibility and decoupling, it is recommended to work with external storage using PV and PVC
- Persistent Volume (PV) is a non-namespaced Kubernetes object that connect to different types of external storage
- Persistent Volume Claim (PVC) contains a specification of required storage needs
- The Pod Volume Spec refers to the PVC to connect to a specific type of storage
- Using this approach guarantees that the developer doesn't have to know anything about storage specifics

Using Dynamic Storage Provisioning

- PVs can be created manually
- Alternatively, PVs can be dynamically provisioned
- To dynamically provision a PV, a StorageClass must be referred to in the PVC
- Storage classes are using classes that are defined by the administrator
 - It's up to the admin to decide what to use to categorize: storage type, backup type, QoS, or something else
- Storage classes are also using provisioners that connect to a specific storage type
 - Each provisioner has its own parameters
- A default storage class can be used, alternatively storage classes can be defined manually

Understanding Default StorageClass

- A Default StorageClass can be used to always specify a storage type so that it doesn't have to be defined in the PVC spec anymore
- If no specific storage class is defined, the default storage class is used
- Often, the hostPath storage type is used for this purpose
- Notice that hostPath should not be used in clusters, as continuous access to storage cannot be guaranteed!
- Using the Default StorageClass provided by a cluster is simple: nothing has to be specified in the PVC definition

Demo: Using Default StorageClass

- A default storage class is provided for each cluster type
- On Minikube, it is set to hostPath
- The example YAML file `default-storageclass.yaml` shows how the default storageclass can be used

Understanding ConfigMaps

- The goal of the ConfigMap is to separate configuration from code
- ConfigMaps are clear-text
- They can be used in three different ways:
 - Make variables available within a Pod
 - Provide command line arguments
 - Mount them on a location where the application expects to find a configuration file
- Secrets are encoded ConfigMaps which can be used to store sensitive data
- ConfigMaps must be created before the pods that are using them

Understanding ConfigMap Sources

- ConfigMaps can be created from different sources
 - Directories: uses multiple files in a directory
 - Files: puts the contents of a file in the ConfigMap
 - Literal Values: useful to provide variables and command arguments that are to be used by a Pod

Procedure Overview: Creating ConfigMaps

- Start by defining the ConfigMap and create it
 - Consider the different sources that can be used
 - **kubectl create cm myconf --from-file=my.conf**
 - **kubectl create cm variables --from-env-file=variables**
 - **kubectl create cm special --from-literal=VAR3=cow --from-literal=VAR4=goat**
 - Verify creation, using **kubectl describe cm <cmname>**
- Use **--from-file** to put the contents of a config file in the ConfigMap
- Use **--from-env-file** to define variables
- Use **--from-literal** to define variables or command line arguments

Procedure Overview: Using ConfigMaps

- To include *variables* from a ConfigMap:
envFrom:
 - configMapRef:
name: ConfigMapName
- To include *config files* from a ConfigMap:
volumes:
 - configMap:
name: configMapName
items:
 - key: my-custom.conf
path: default.conf

Demo: Creating a ConfigMap from a File

Files are provided in the GitHub repository

- Consider the contents of the file **variables**
- Create the ConfigMap: **kubectl create cm variables --from-env-file=variables**
- Verify creation: **kubectl describe cm variables**
- Create a Pod: **kubectl create -f cm-vars.yaml**
- Check that the variables are available: **kubectl logs po/cm-vars**

Demo: Creating a ConfigMap from a Literal

- **kubectrl create cm cm-literal --from-literal=VAR3=goat
--from-literal=VAR4=cow**
- **kubectrl get cm cm-literal**

Demo: Using ConfigMaps for

- **cat nginx-custom-config.conf**
- **kubectl create cm nginx-cm --from-file nginx-custom-config.conf**
- **kubectl get cm nginx-cm -o yaml**
- **kubectl create -f nginx-cm.yaml**
- **kubectl exec -it nginx-cm /bin/bash**
 - **cat /etc/nginx/conf.d/default.conf**

Demo: Creating Secrets

- **kubectI create secret generic my-secret**
--from-file=ssh-privatekey=/home/student/.ssh/id_rsa
--from-literal=passphrase=password
- **kubectI describe secret my-secret**

Understanding Services

- A service is an API object that is used to expose a logical set of Pods
- It is also an abstraction that defines the logical set of Pods and a policy how to access them
- By working this way, Kubernetes service resource implement microservices
- The set of Pods that is targeted by a service is determined by a selector (which is a label)
- The Kubernetes cluster controller will continuously scan for pods that match the selector and include these in the service

Understanding Services Decoupling

- Services are independent resources, they are not directly connected to a deployment
- The only thing they do is watch for a deployment (or other resources) based on the selector label that is specified in the service
- This allows for services to connect to multiple deployments, and if that happens, Kubernetes will automatically load balance between these deployments
- Users connect to services on a fixed IP address / port combination, and the service load balances the connection to one of the endpoint pods it services

Understanding Service Types

- ClusterIP: the default service type. As cluster IP is not reachable from the outside, it provides internal access only
- NodePort: allocates a port on the kubelet nodes to expose the service externally
- Loadbalancer: implements a load balancer for direct external access. Only available in public cloud
- ExternalName: a service object that works with DNS names

Understanding the ClusterIP type in Microservices

- The ClusterIP service type doesn't seem to be very useful, as it doesn't allow access to a service from the outside
- For communication of applications in a microservice architecture, it is perfect though: it does expose the application on the cluster IP network, but makes it inaccessible for external users

Understanding Service Ports

While working with services, different Ports are specified

- targetport: this is the port on the application that the service object connects to
- port: this is what maps from the cluster nodes to the targetport in the application to make the port forwarding connection
- nodeport: this is what is exposed externally

Demo: Exposing Applications Using Services

- **kubectl create deployment nginxsvc --image=nginx**
- **kubectl scale deployment nginxsvc --replicas=3**
- **kubectl expose deployment nginxsvc --port=80**
- **kubectl describe svc nginxsvc # look for endpoints**
- **kubectl get svc nginx -o=yaml**
- **kubectl get svc**
- **kubectl get endpoints**

Demo: Accessing Deployments by Services

- `minikube ssh`
- `curl http://svc-ip-address`
- `exit`
- `kubectl edit svc nginxsvc`
 - ...
 - `protocol: TCP`
 - `nodePort: 32000`
 - `type: NodePort`
- `kubectl get svc`
- (from host): `curl http://$(minikube ip):32000`

Understanding Ingress

- Ingress exposes HTTP and HTTPS routes from outside the cluster to services within the cluster
- Traffic routing is controlled by rules defined on the Ingress resource
- Ingress can be configured to do the following
 - Give services externally reachable URLs
 - Load balance traffic
 - Terminate SSL/TLS
 - Offer name based virtual hosting
- An ingress controller is used to realize Ingress
 - On Minikube, use **minikube addons enable ingress**

Demo: Creating Ingress - 1

Note: this demo continues on the demo in lesson 13.3

- **minikube addons enable ingress**
- **kubectl get deployment**
- **kubectl get svc nginxsvc**
- **curl http://\$(minikube ip):32000**
- **vim nginxsvc-ingress.yaml**
- **kubectl apply -f nginxsvc-ingress.yaml**
- **kubectl get ingress**
- **sudo vim /etc/hosts**
 - **\$(minikube ip) nginxsvc.info**
- **curl nginxsvc.info**
- Note this doesn't work!

Demo: Creating Ingress - 2

- Previous demo doesn't work; the point of using Ingress is to have one centralized ingress into the cluster and that doesn't work with the NodePort service type
- **kubectl edit svc nginxsvc**
 - remove **nodePort: 32000**
 - change to **type: ClusterIP**
- **vim nginxsvc-ingress.yaml**
 - change to **servicePort: 80**
- **kubectl apply -f nginxsvc-ingress.yaml**
- **curl nginxsvc.info**

Understanding Kustomize

- **kustomize** is a Kubernetes feature, included since 1.14, that uses a file with the name **kustomization.yaml** to store instructions on the changes a user wants to make to a set of resources
- This is convenient for applying changes to input files that the user does not control himself, and which contents may change because of new versions appearing in Git
- So it's a way to modify existing configuration files, and thus implement changes in a declarative way
- Use **kubectl apply -k ./** in the directory with the **kustomization.yaml** and the files it refers to to apply changes
- Use **kubectl delete -k ./** in the same directory to delete all that was created by the Kustomization

Understanding a Sample kustomization.yaml

```
resources:          # defines which resources (in YAML files) apply
- deployment.yaml
- service.yaml
namePrefix: test-    # specifies a prefix should be added to all names
namespace: testing   # objects will be created in this specific namespace
commonLabels:        # labels that will be added to all objects
environment: testing
```

Using Kustomization Overlays - 1

- Kustomization can be used to define a base configuration, as well as multiple deployment scenarios (overlays) as in dev, staging and prod for instance
- In such a configuration, the main kustomization.yaml defines the structure:
 - base
 - deployment.yaml
 - service.yaml
 - kustomization.yaml
 - overlays
 - dev
 - kustomization.yaml
 - staging
 - kustomization.yaml
 - prod
 - kustomization.yaml

Using Kustomization Overlays - 2

- In each of the overlays/{dev,staging.prod}/kustomization.yaml, users would reference the base configuration in the resources field, and specify changes for that specific environment:

resources:

- ../../base

namePrefix: dev-

namespace: development

commonLabels:

environment: development

Demo: Using Kustomization.yaml

- **cat deployment.yaml**
- **cat service.yaml**
- **kubectl apply -f deployment.yaml service.yaml**
- **cat kustomization.yaml**
- **kubectl apply -k .**

Poll Question 6

Which of the following topics did not get enough coverage (choose multiple)

- Understanding Microservices
- Using Git in DevOps
- Managing Docker Overview
- The Role of Containers in Orchestration
- Understanding Kubernetes
- Creating Container Based Microservices

Poll Question 7

Which of the following topics did you like most in this course (choose multiple)

- Understanding Microservices
- Using Git in DevOps
- Managing Docker Overview
- The Role of Containers in Orchestration
- Understanding Kubernetes
- Creating Container Based Microservices

Further Learning: Live Courses

- Containers in 4 Hours: Docker and Podman
- Kubernetes in 4 Hours
- Getting Started with OpenShift
- Certified Kubernetes Application Developer (CKAD) Crash Course
- Certified Kubernetes Administrator (CKA) Crash Course

Further Learning: Recorded Courses

- Hands-on Kubernetes
- Getting Started with Kubernetes LiveLessons 2nd Edition
- Red Hat OpenShift Fundamentals 3/
- Certified Kubernetes Application Developer (CKAD)
- Certified Kubernetes Administrator (CKA)