



Microservices & Apache Kafka®

Part 2 – Building Event-Driven Services with Apache Kafka

Series Schedule



- Session 1: The Data Dichotomy: Rethinking the way we treat data and services
- **Session 2: Building Event-Driven Services with Apache Kafka**
- Session 3: Putting the 'Micro' into Microservices with Stateful Stream Processing



In this talk

1. How we traditionally build services with REST
2. How Event Driven Services are different
3. Building an Immutable, Shared Narrative.
4. Leveraging Materialized Views.
5. Relating to CQRS & Event Sourcing
6. Pulling it all together

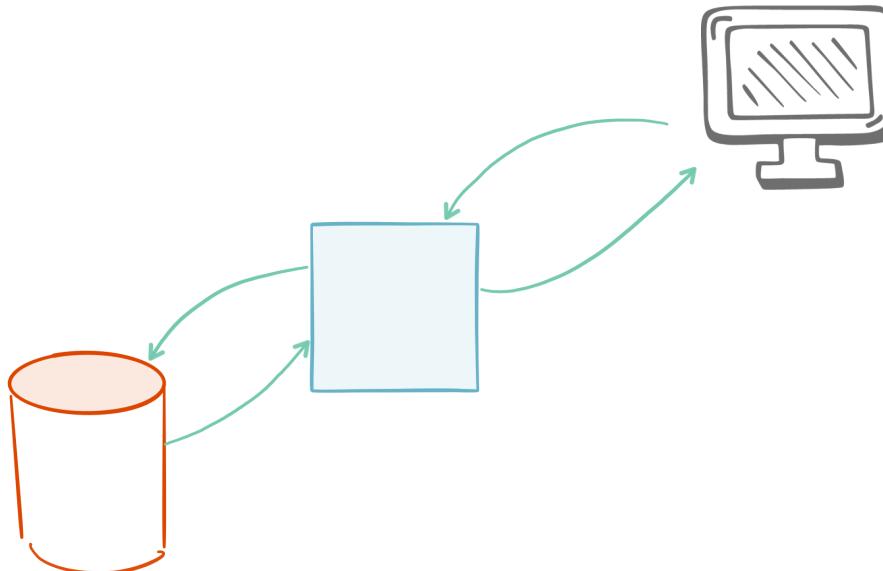


Companies evolve in different ways

- Internet companies
- Enterprise companies

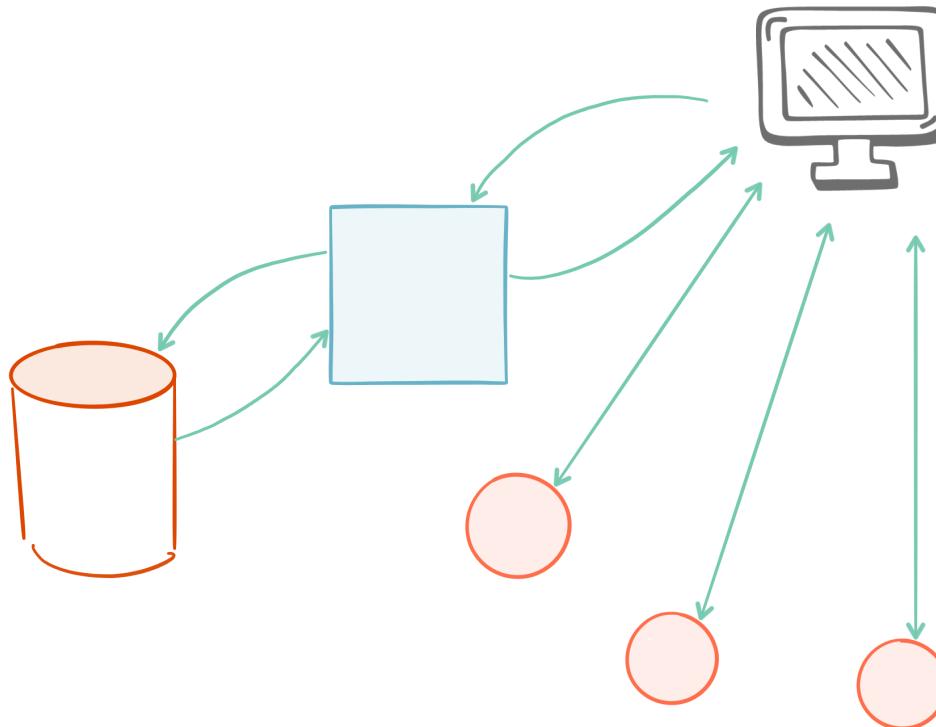


Internet companies typically grow
from front-facing websites.



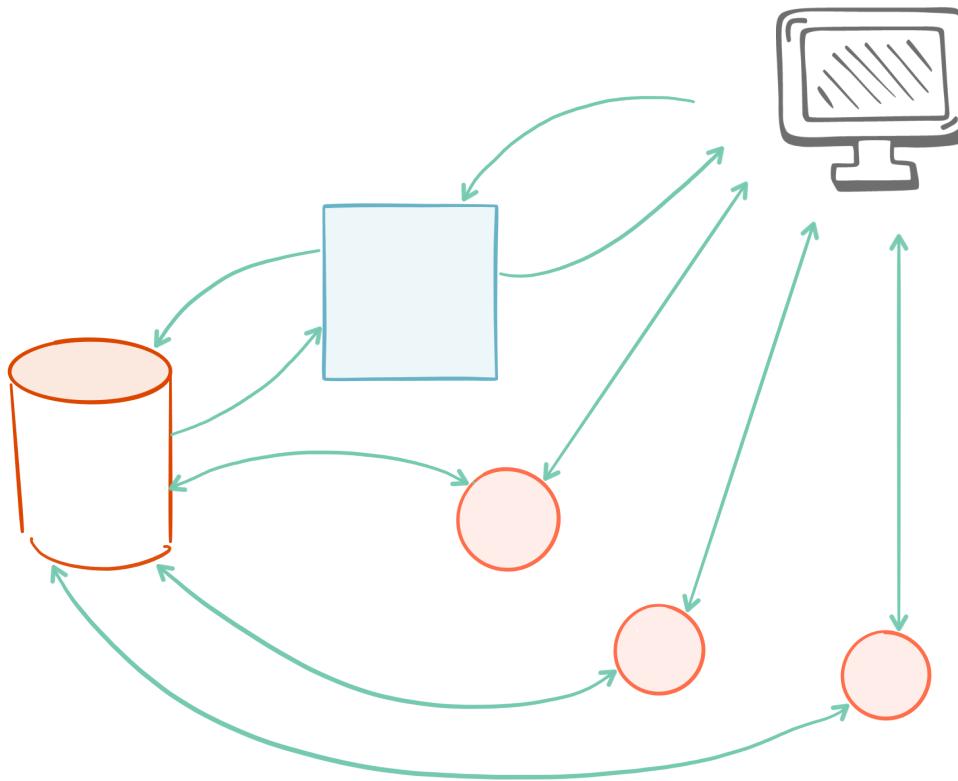
Growing fast

Carve pieces off

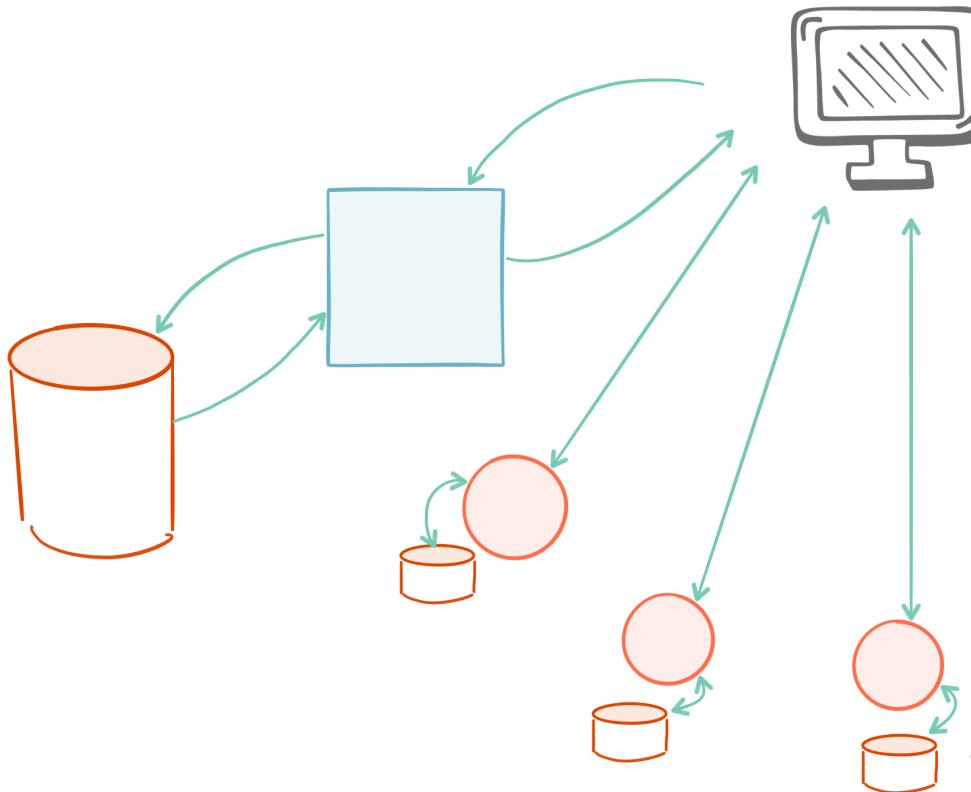




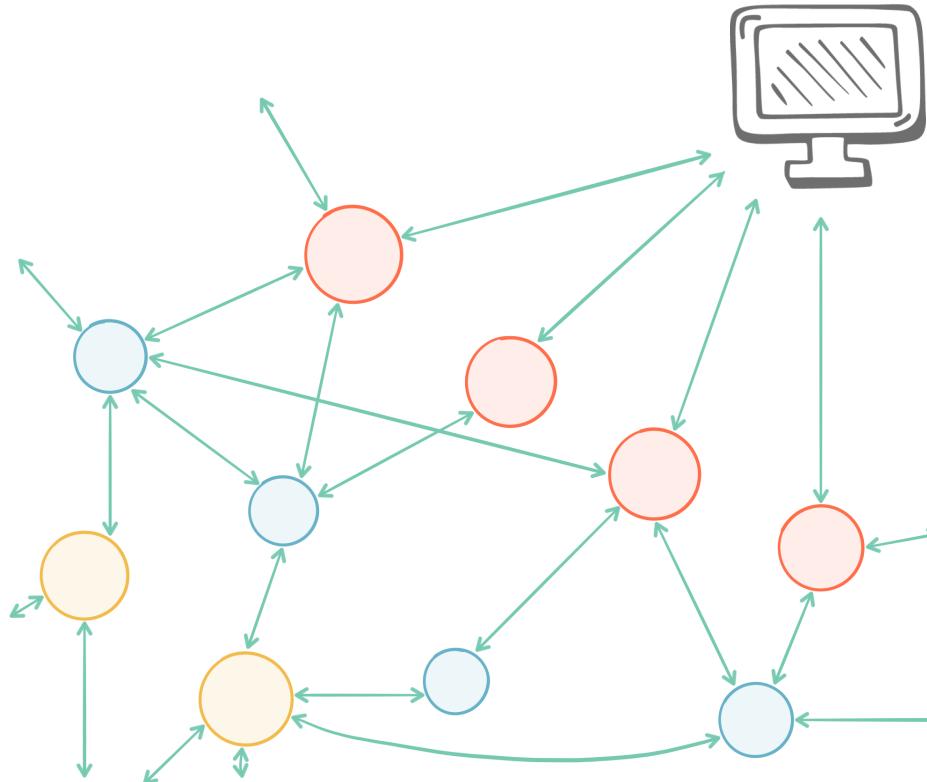
If stateful, they might share a DB



Or they might use their own



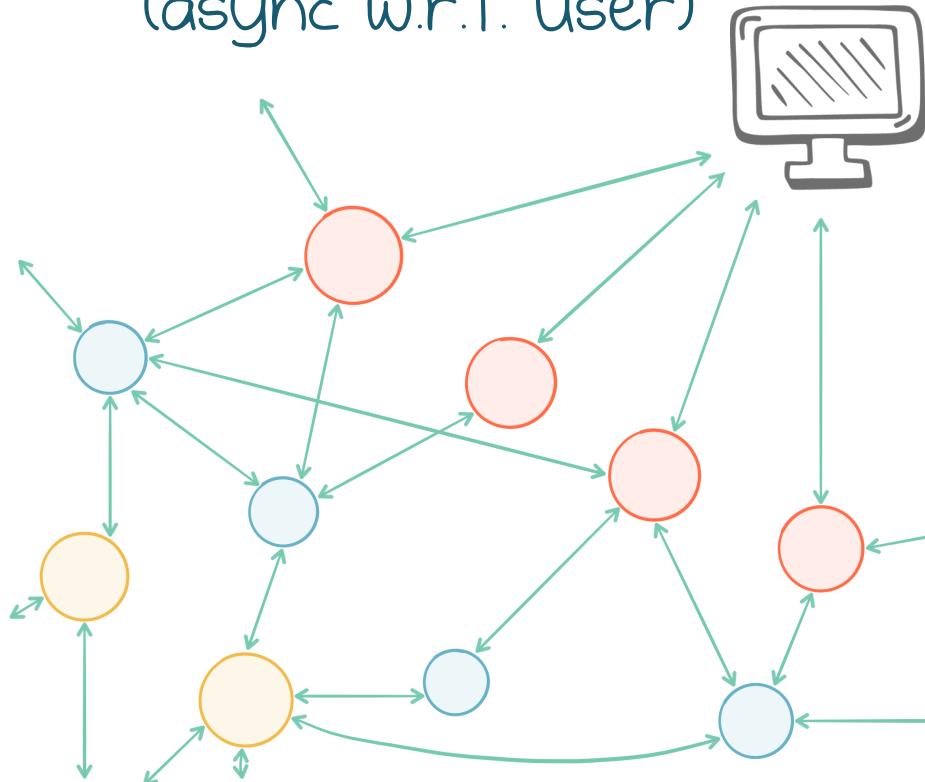
Over time => more complex



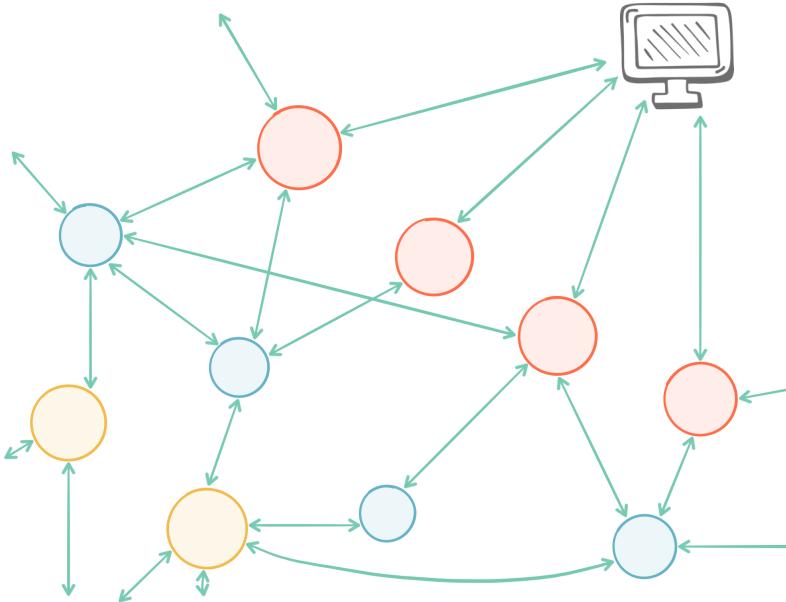
Over time => more async



(async w.r.t. user)



Chained blocking commands



Buffering, handling failure, backpressure, scaling etc all get pushed into service's responsibilities



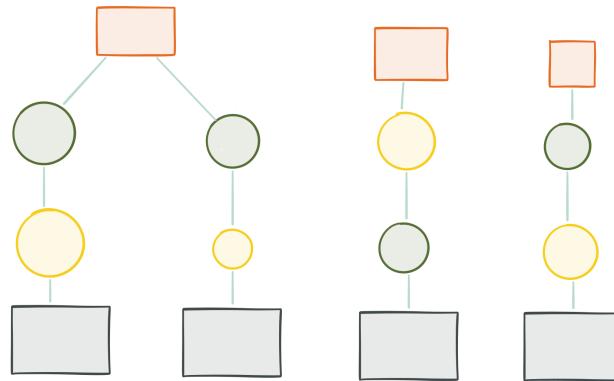
The synchronous world of
request response protocols
leads to tight, point-to-point
couplings.



Enterprises are typically different



Enterprise companies form as heterogeneous application conglomerates, often in silos



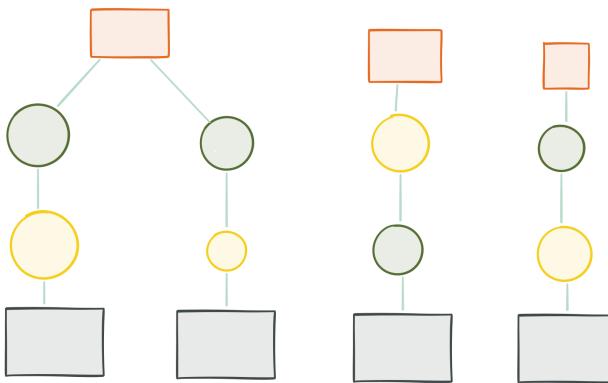


The start life as independent islands



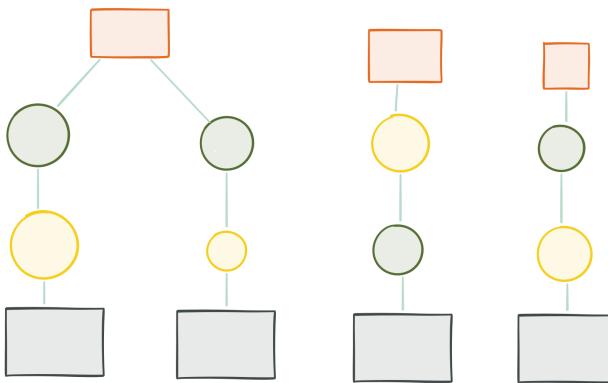


Linked together by files or via messaging

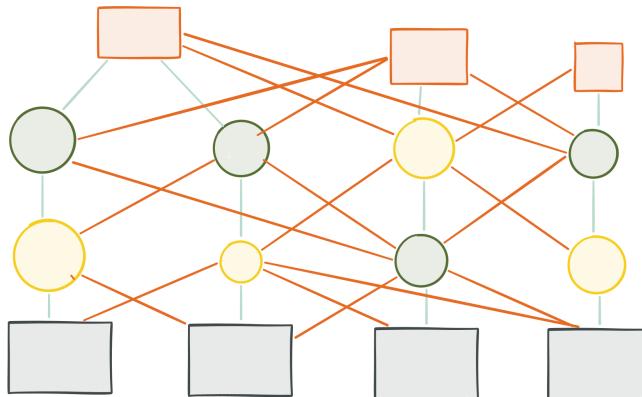




Interactions are async from day-1



Forwarding of data gets messy: Telephone game

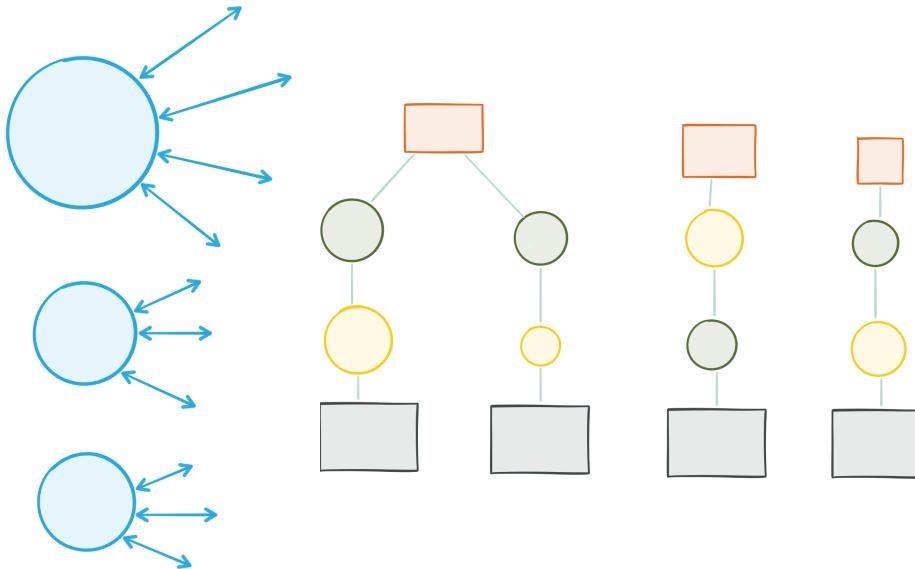




Both these approaches often end up at
the same point, but for different
reasons



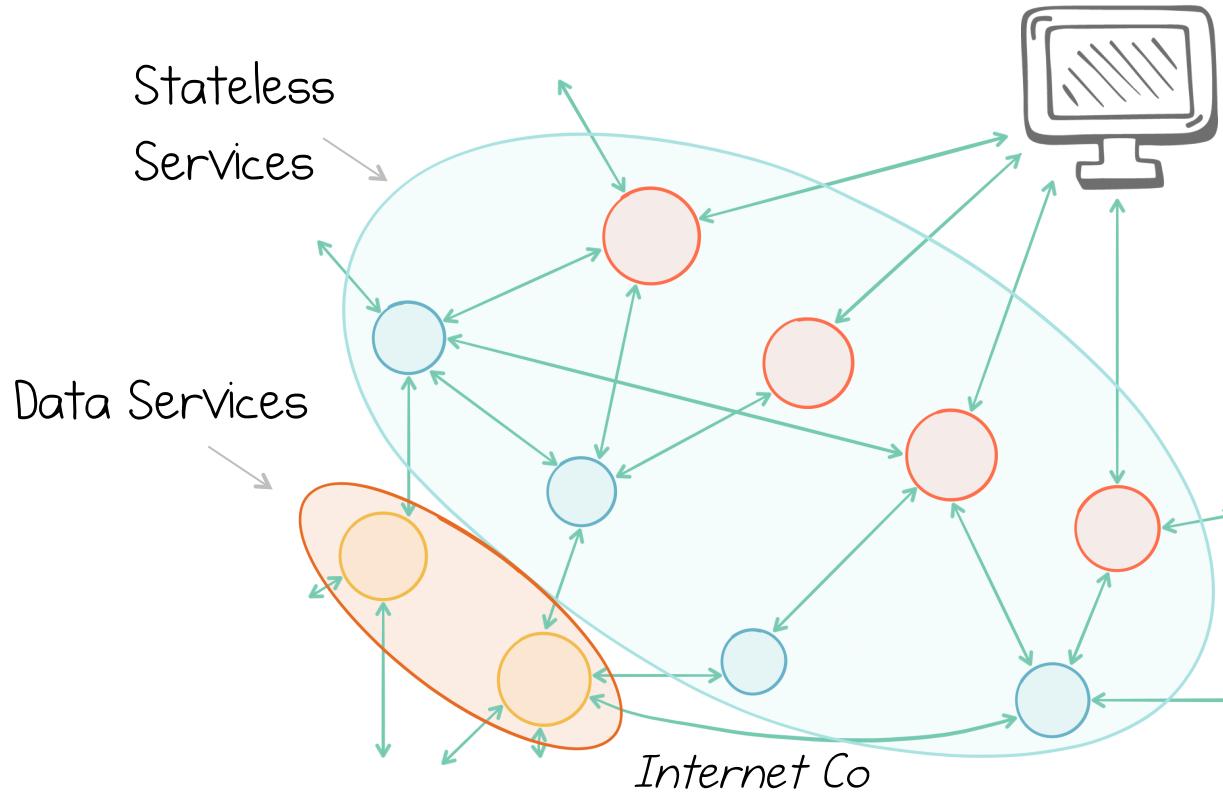
Data is herded into specific services



Enterprise



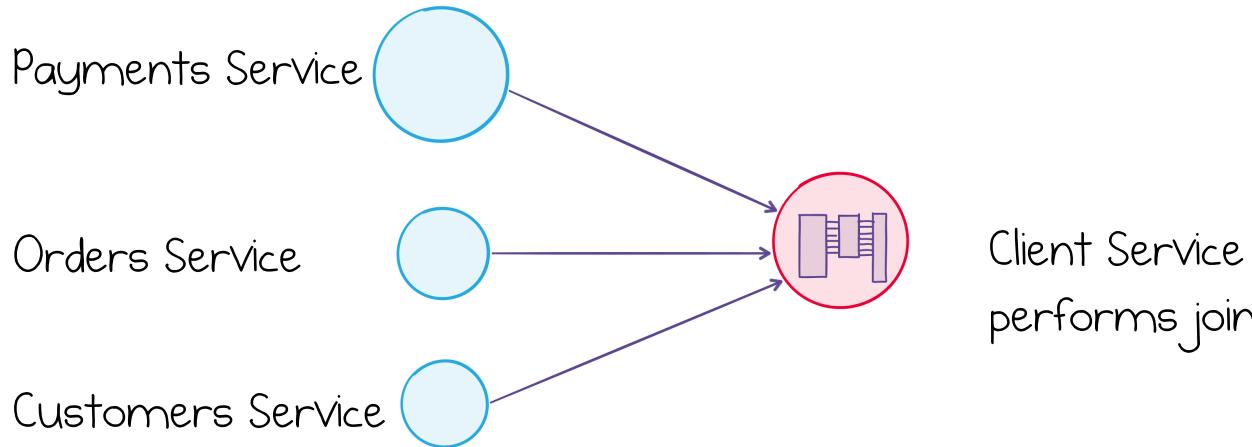
Data is herded into specific services





Lots of Data Services

=> Distributed Join Problem

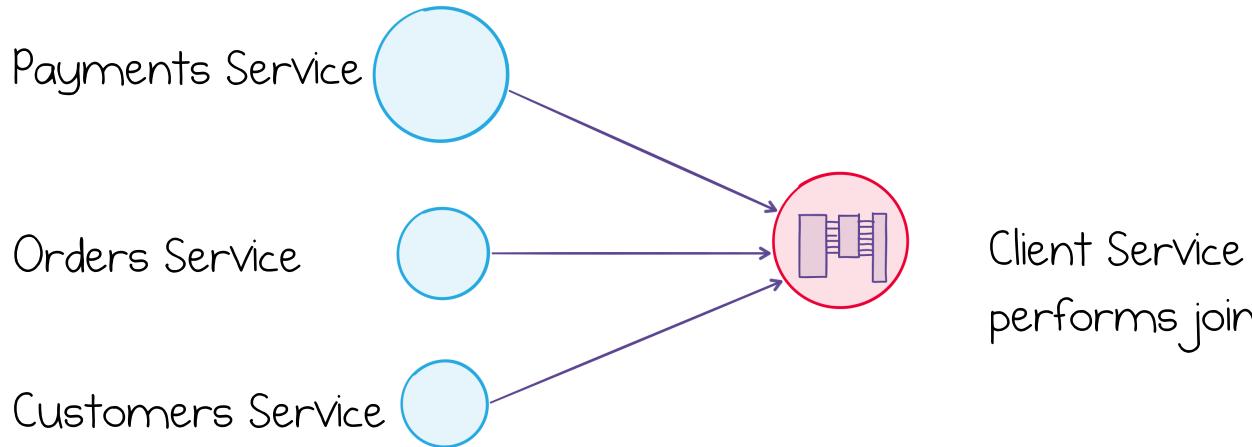


This is a problem databases find hard,
and they're highly tuned for it!



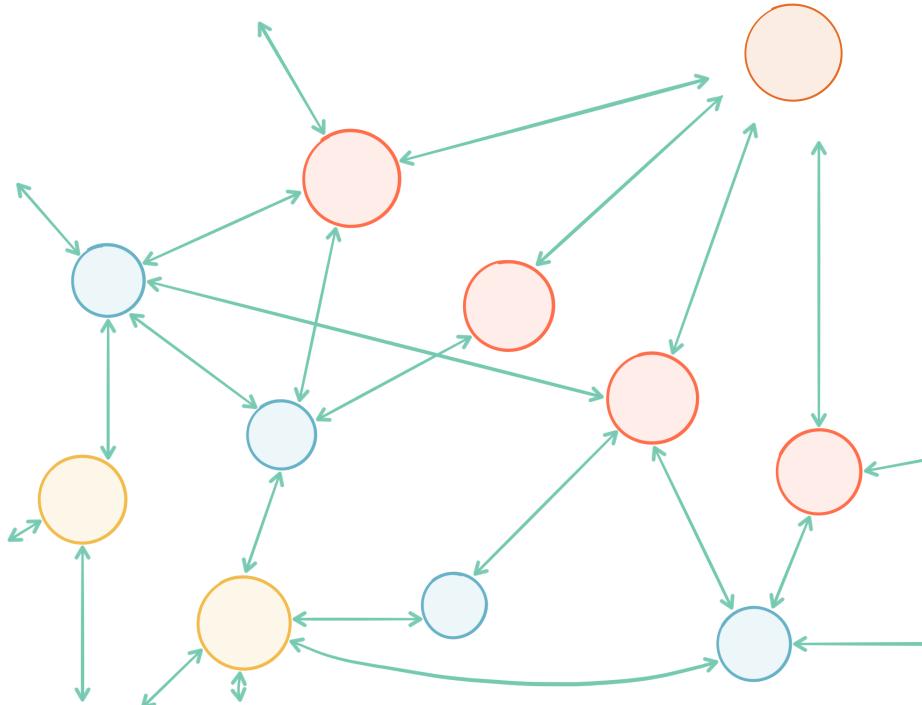
Lots of Data Services

=> Distributed Join Problem



TESTING BECOMES DIFFICULT!

In short: the high degree of “connectedness” makes it hard to evolve independently and to scale





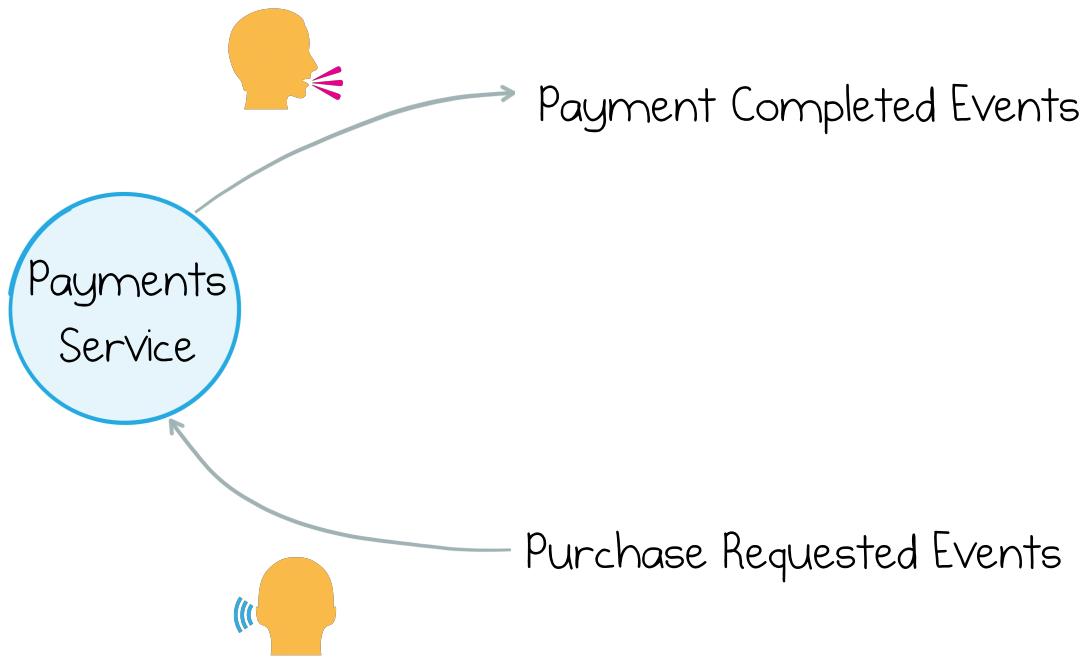
LETS TRY
EVENT DRIVEN!



INTERACT THROUGH
EVENTS.
DON'T TALK TO
SERVICES



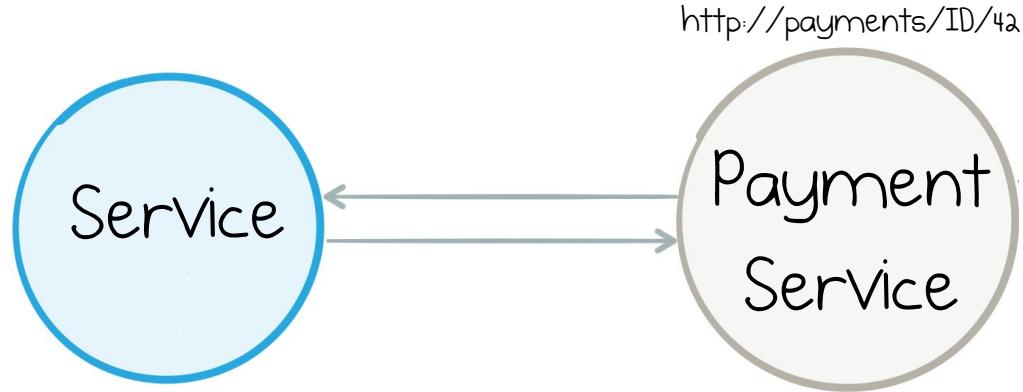
This



Listen & React to Events



Not this



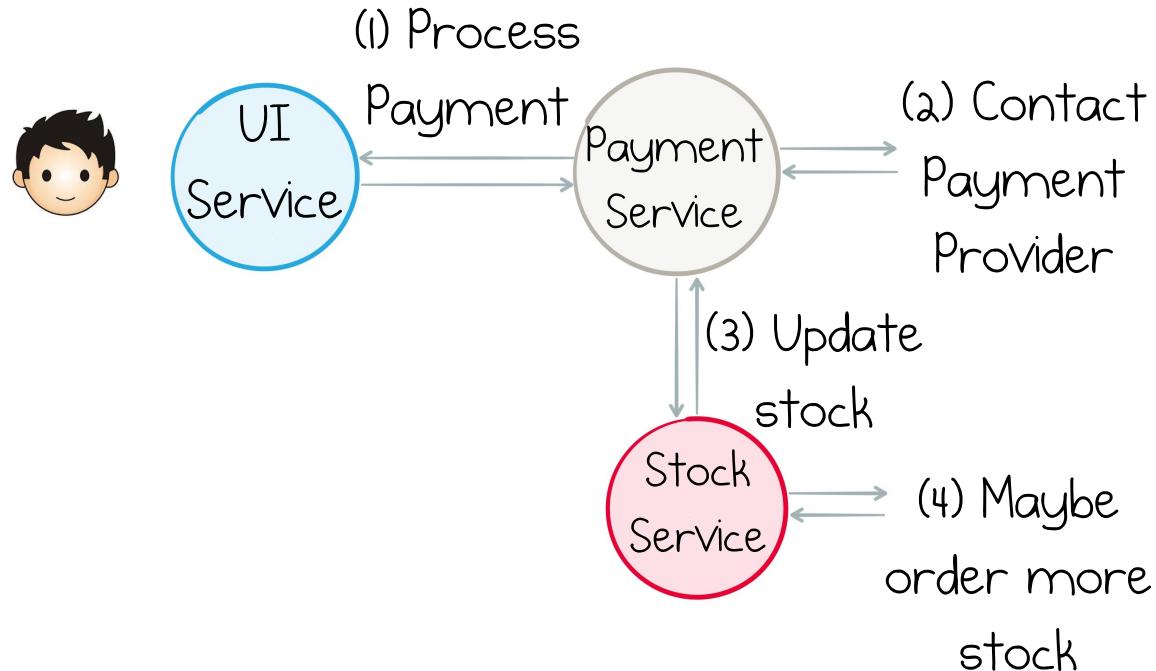
Don't Request/Command from individual services.



Let's take an example...

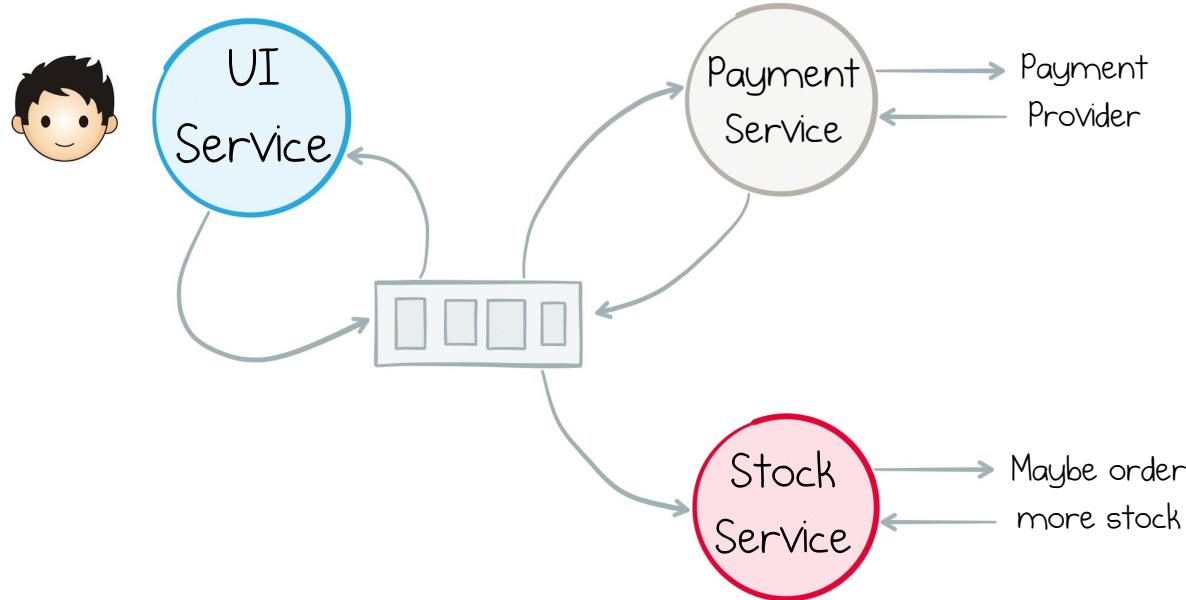


Request Response Example

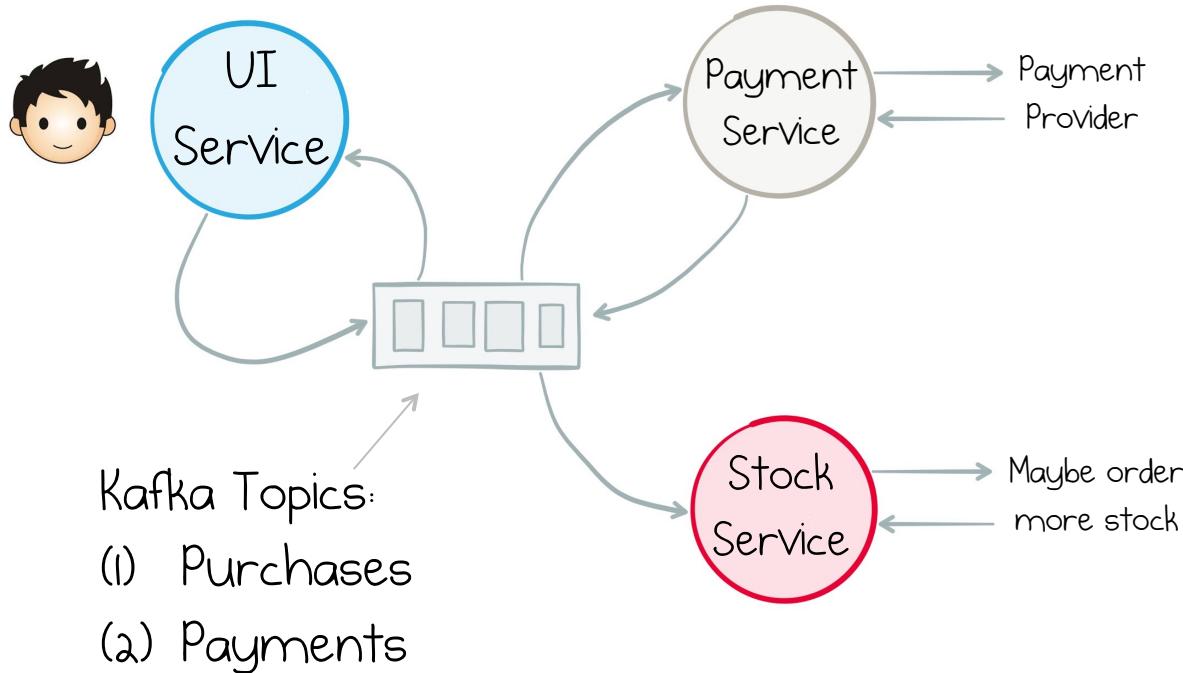




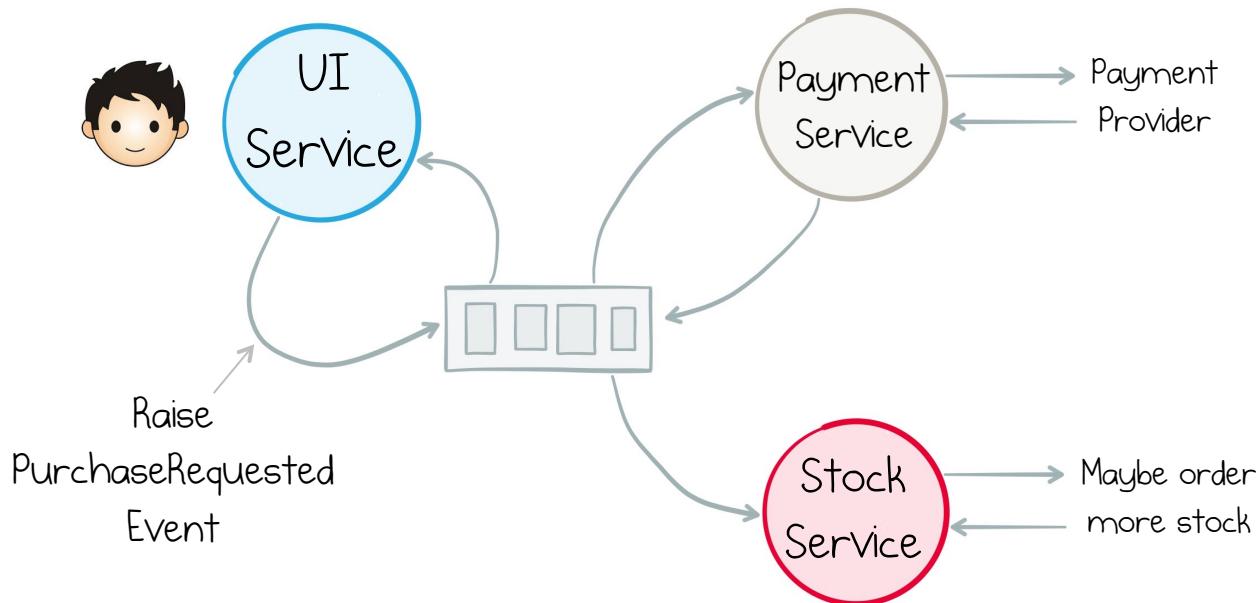
The Event Driven Way



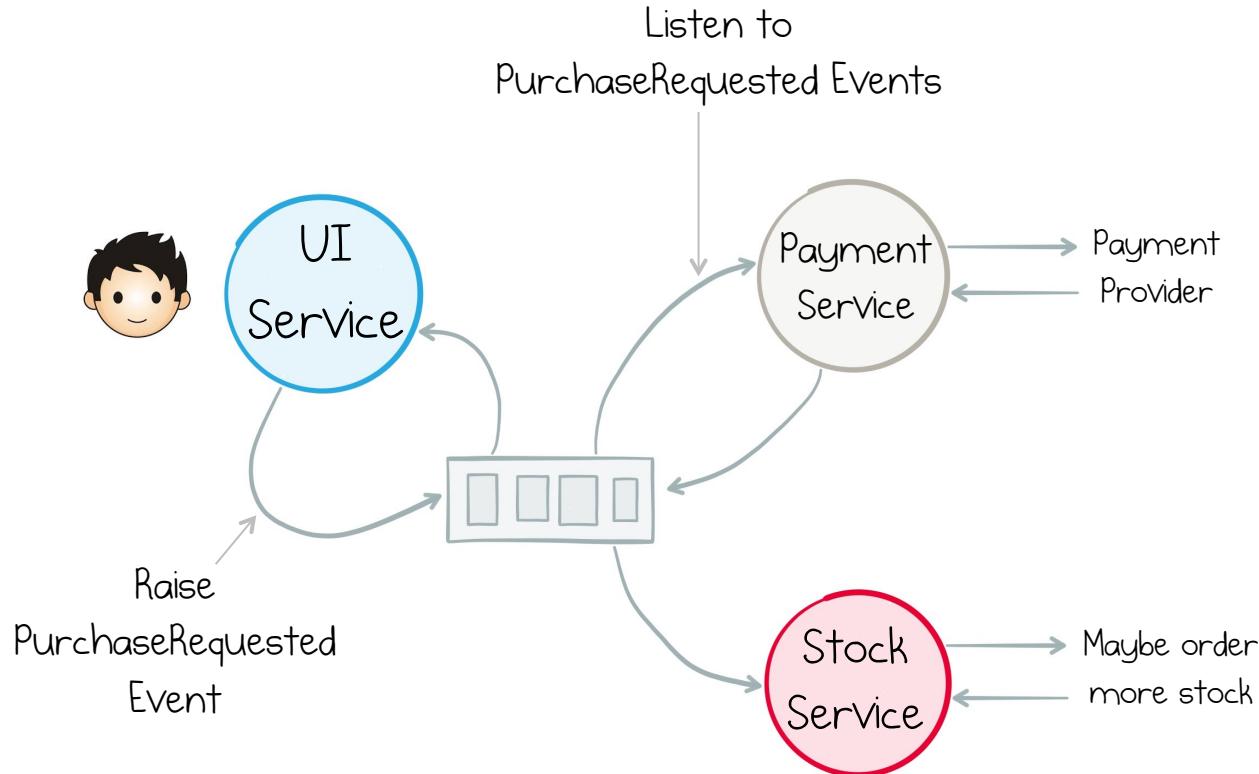
The Event Driven Way



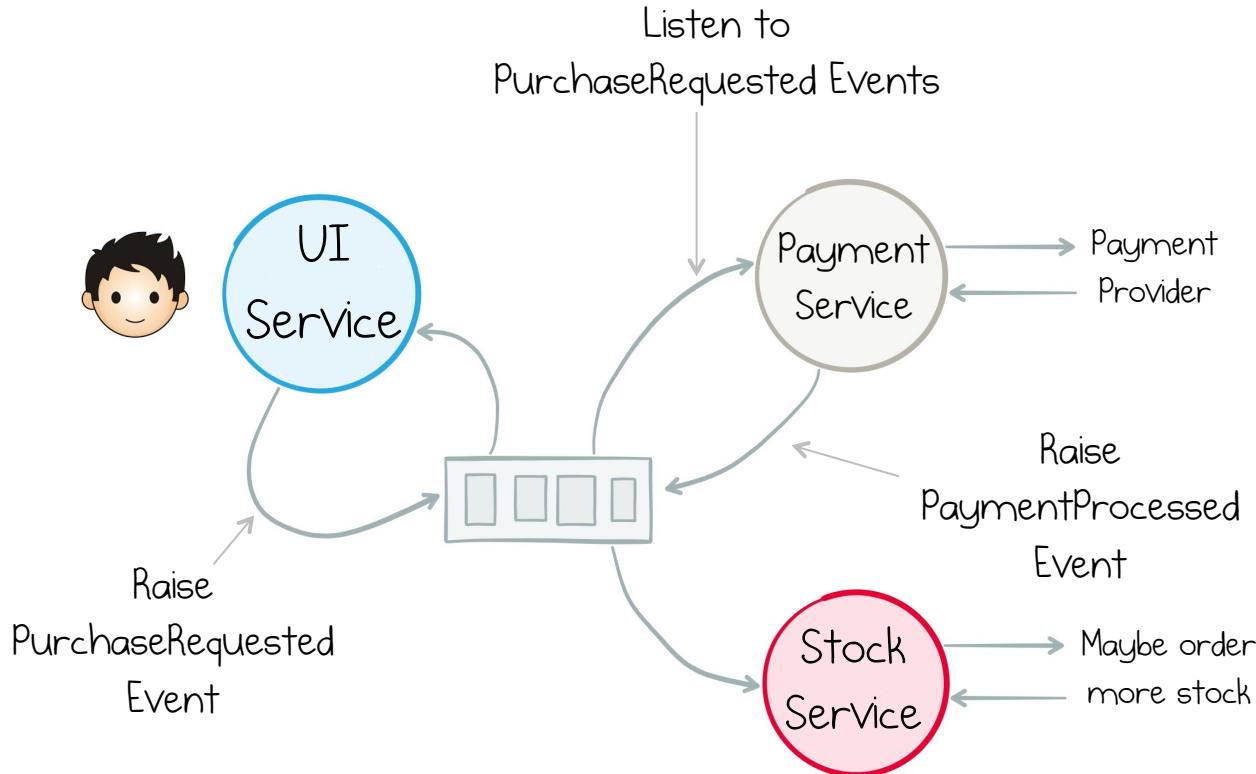
New Order -> PurchaseRequested Event



Payment Service Listens In



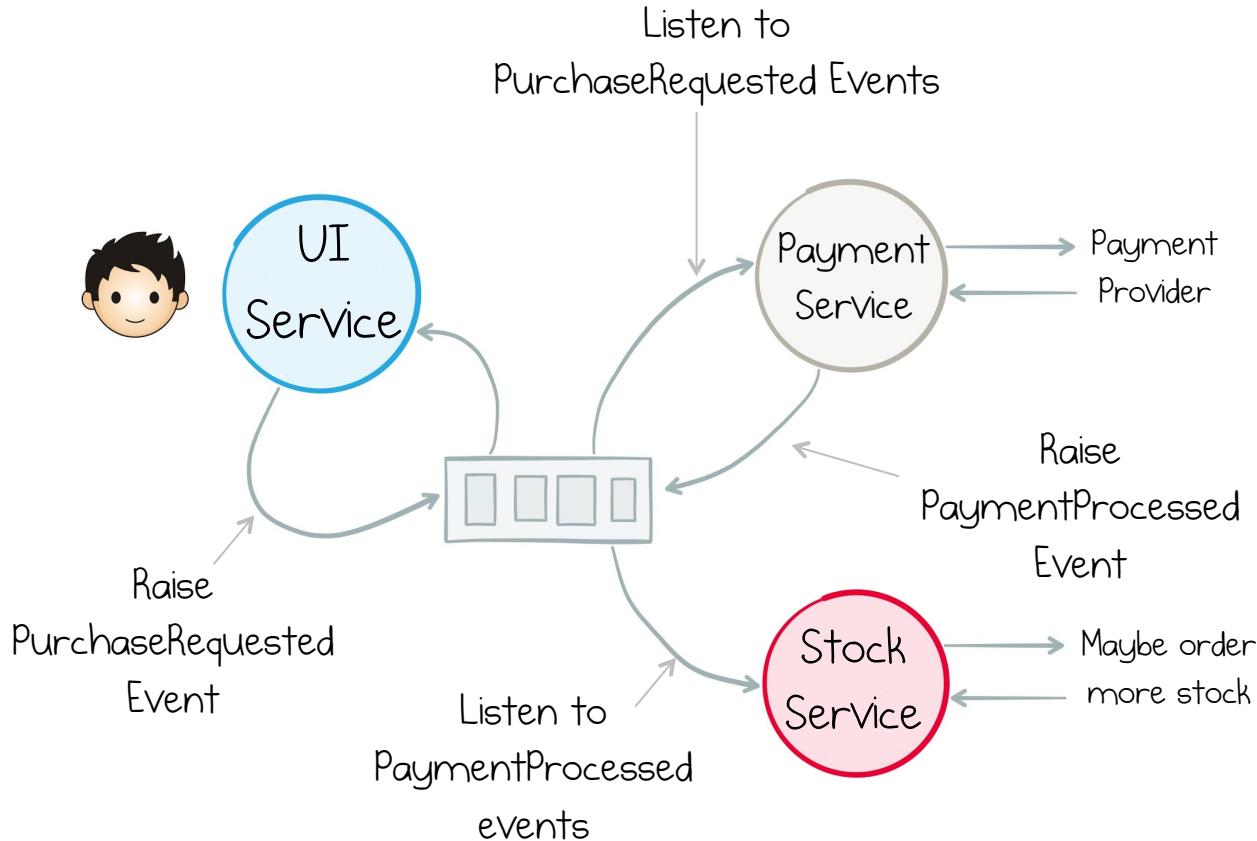
When Complete -> PaymentProcessed



*Some might use the Command/Query Pattern

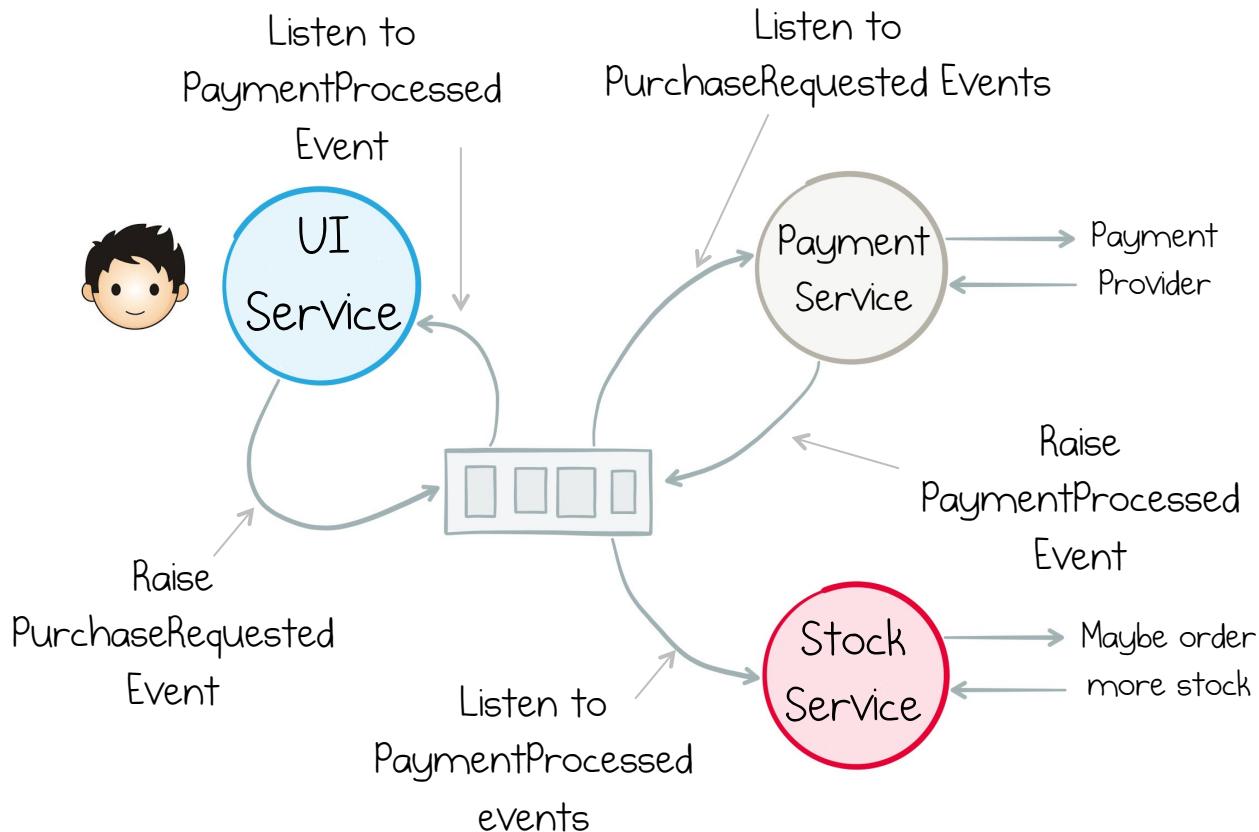


Stock service hooks in too!

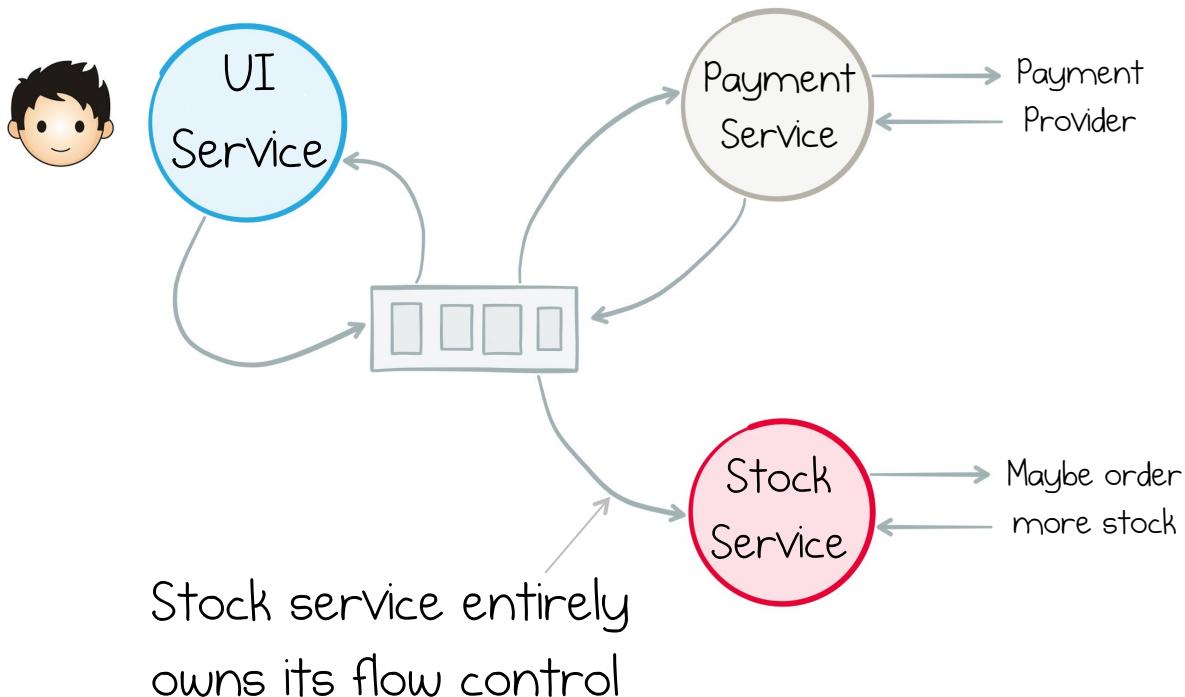




Payment service listens

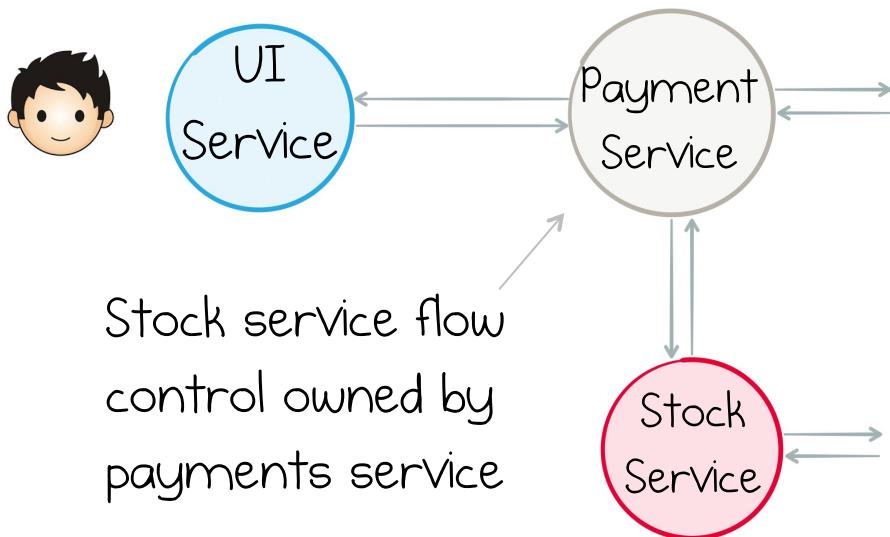


Flow Control is Receiver Driven





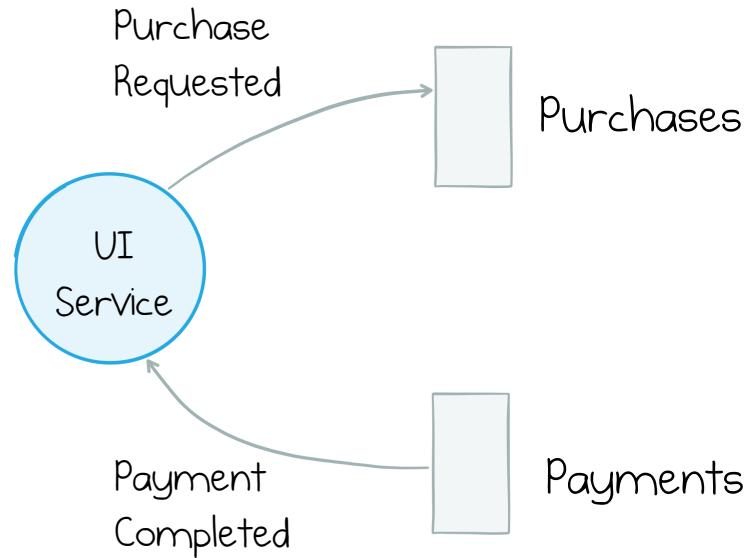
Quite different to Synchronous Model





I. Events are Immutable Facts

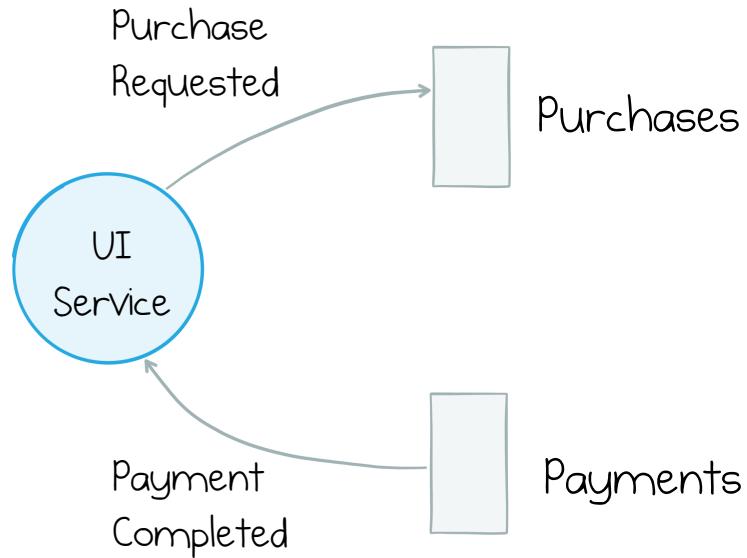
(state changes taken from the real world and journalled)



All the benefits of “Event Sourcing”



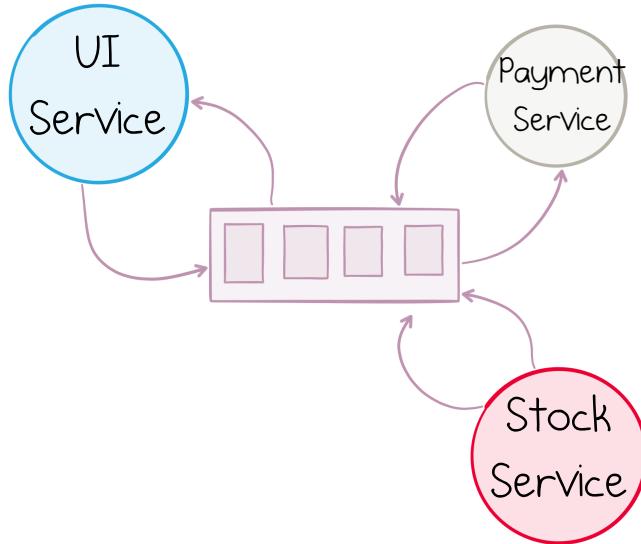
2. Services couple only to data flows



No knowledge of contributing services (RDFC)

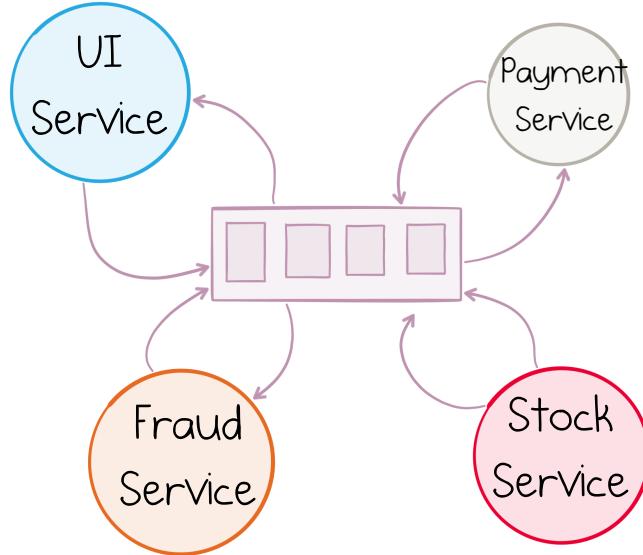


Because coupling is only to Events, not services





The architecture is “Pluggable”

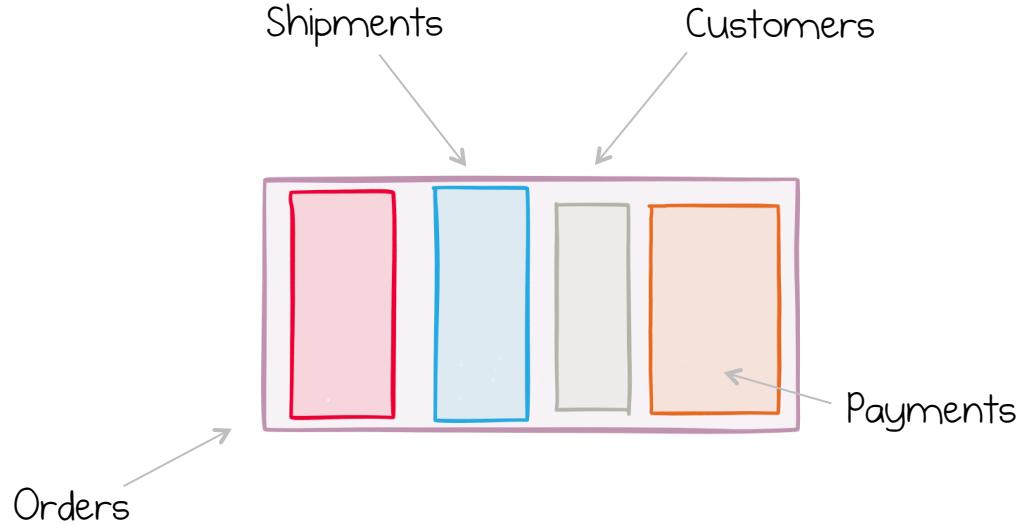




Communication & State Concepts Merge

- The concept of Data & Events become one.
- Kafka is both Event Store and Communication Channel

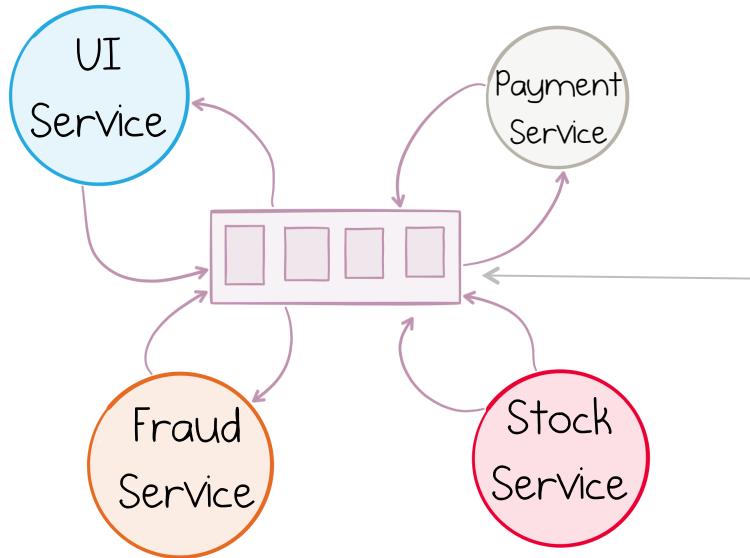
The Events form a Canonical, Shared Narrative



Evolving state of the system over time



Scaling is a concern of the broker, not “upstream” services



Kafka

- Linearly scalable
- Fault Tolerant
- Multi Tenant



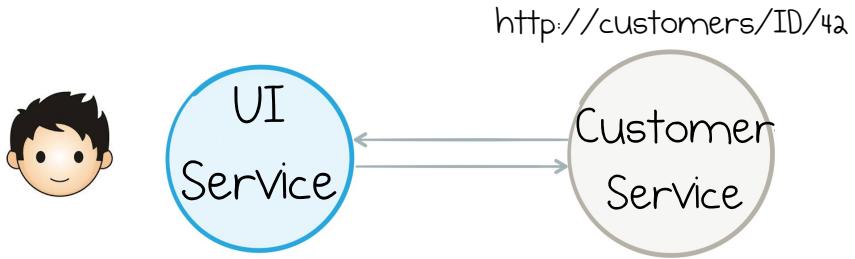
But there is something missing!

How do we look things up?
(i.e. the Q in CQRS)

- What is the address for this customer?
- Is this user allowed to view this stock?
- What is the contents of this user's shopping basket?



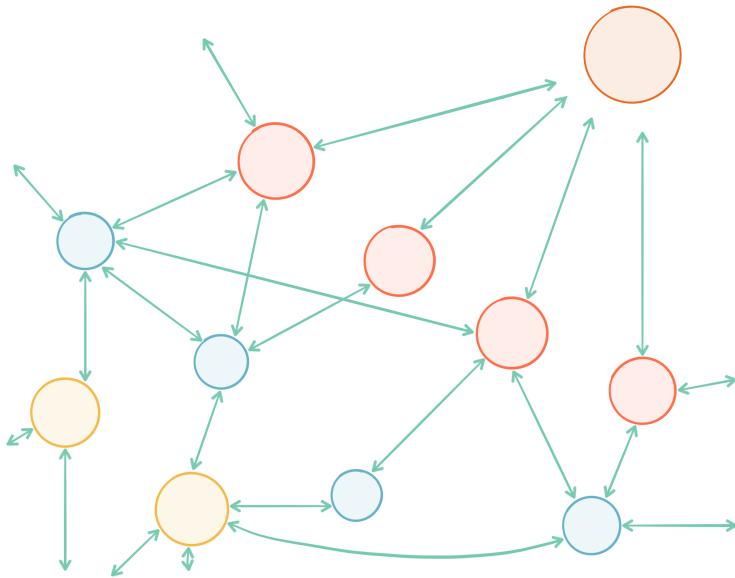
With REST Lookups are natural, if Remote



Go to the relevant service and ask!

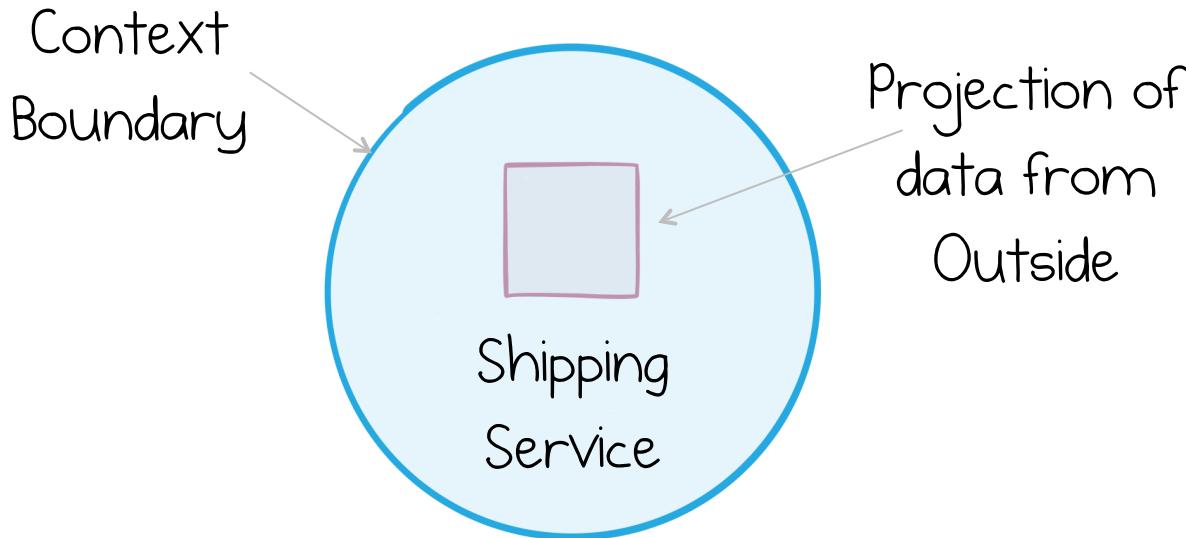


Which creates the dependency problem



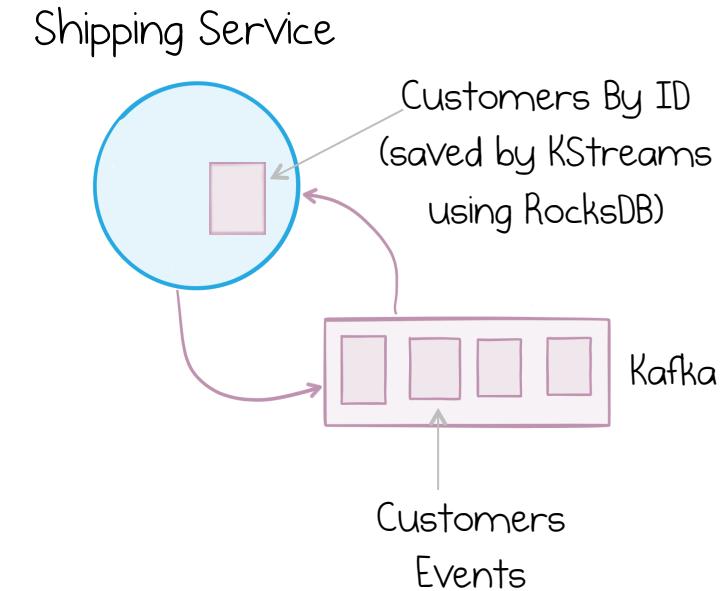
Many services tightly coupled to one another

With Event Driven we create “views” or “projections” inside each bounded context



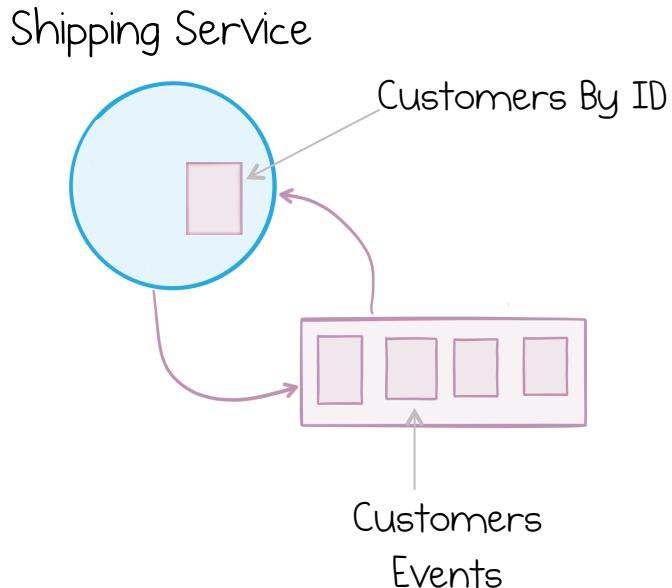


Pattern 1: Local KTables





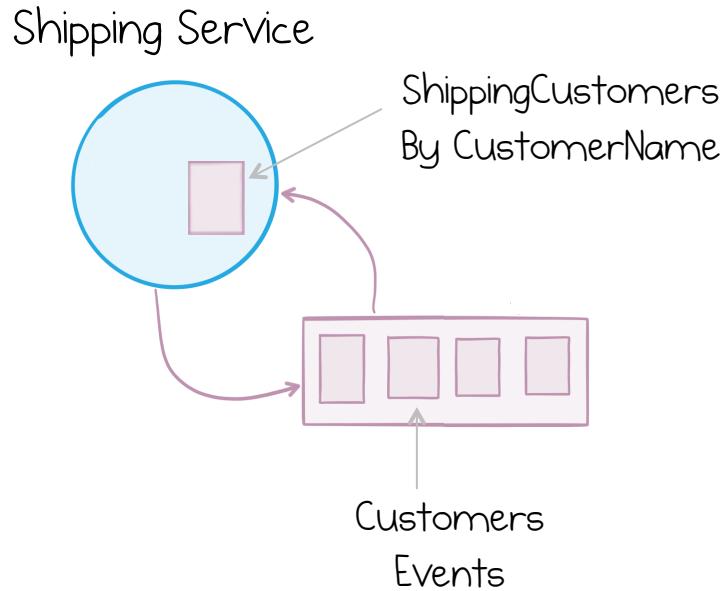
Pattern 1: Local KTables



```
KTable customers =  
builder.table(  
    CustomerId,  
    Customer,  
    "customers-topic",  
    "customer-store");  
  
customers = streams.store(  
    "customers-store"...);  
  
customer = customers.get(42)
```

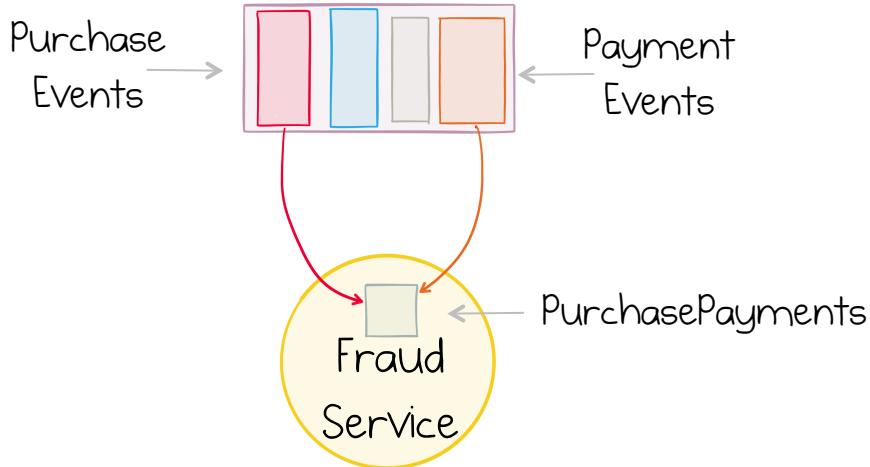


As we're using a streaming engine, we can translate into any Domain Model (projection) we wish





Combine different streams from different services



```
payments.join("purchases", ...)  
.groupByKey()  
.reduce(newValueReducer, "PurchasePaymentsStore")
```

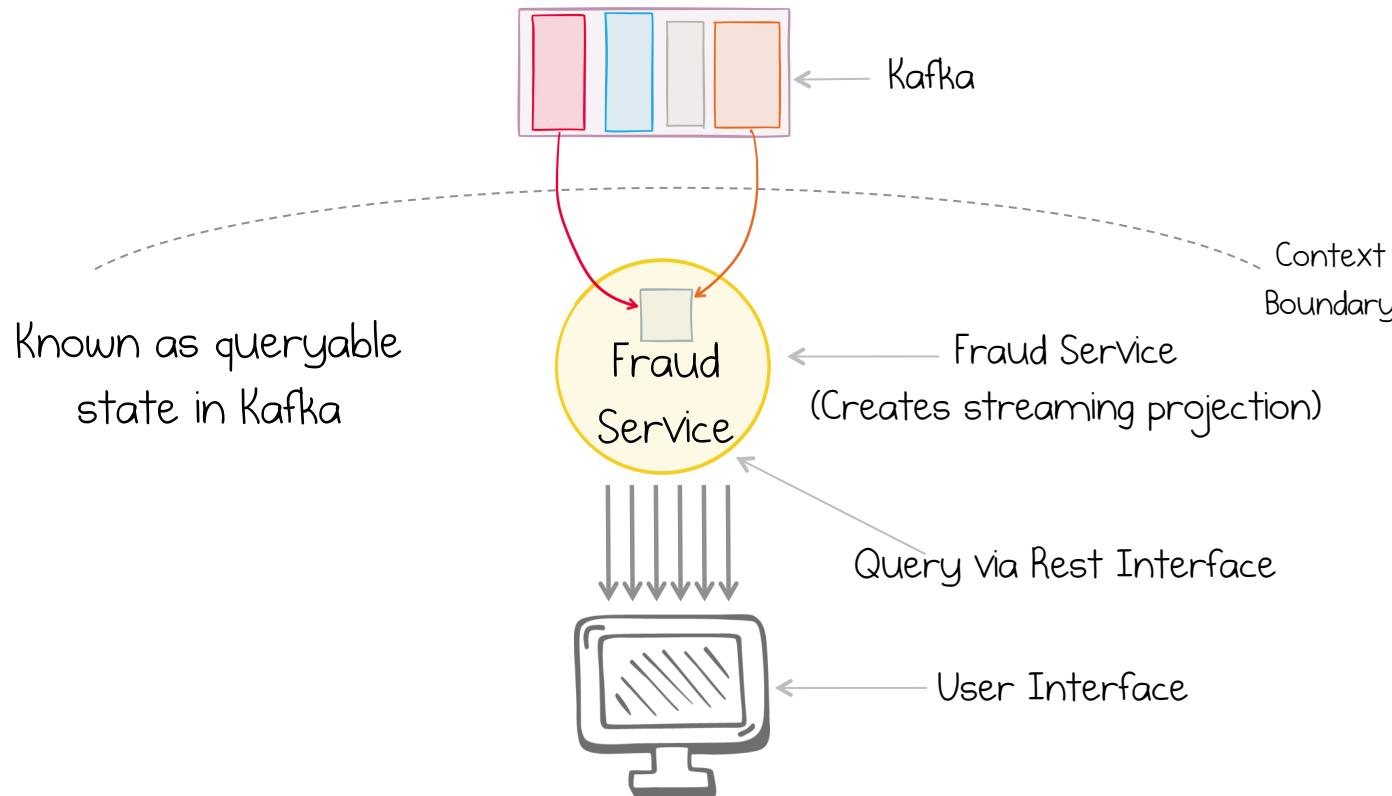


Pattern 1: Local KTables

- Simple to create and query
- Local to each service so fast
- Powerful DSL for transformation
- Inbuilt high availability
- No external database
- Controlled entirely within service's bounded context
- KTable and GlobalKTable allow scale out

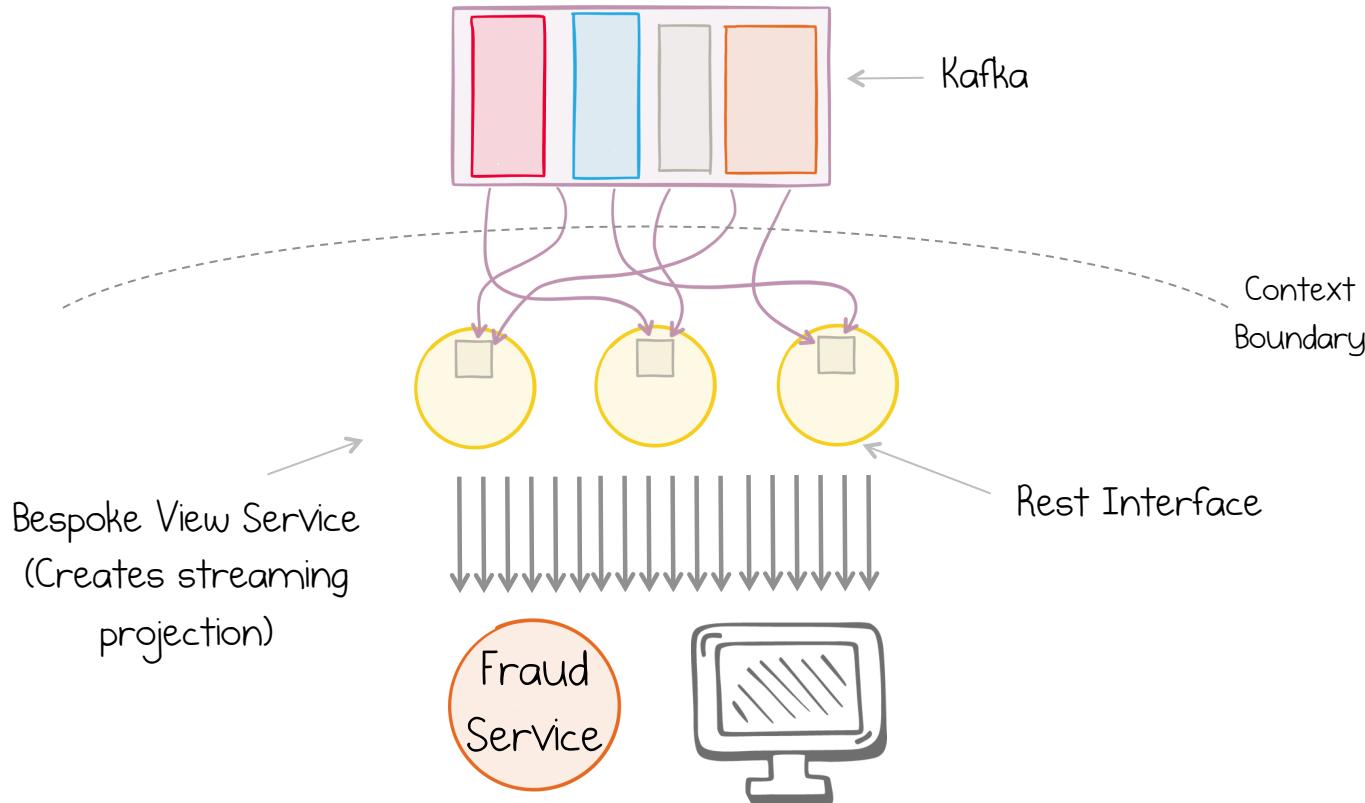


Pattern 2(a): Queryable Interface





Pattern 2(b): Query Service



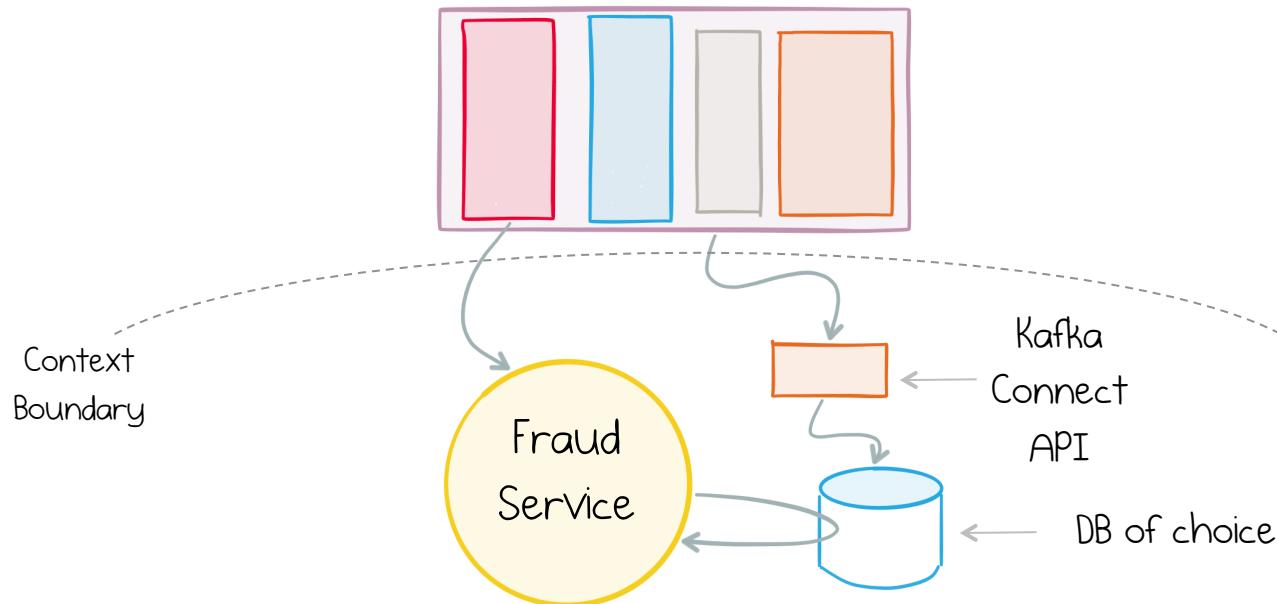
Pattern 2: Exposing Materialised Projections



- Similar to Pattern 1 but with a “Query Interface” (using Kafka’s Queryable State feature)
- Commonly used with UI’s



Pattern 3: External DB



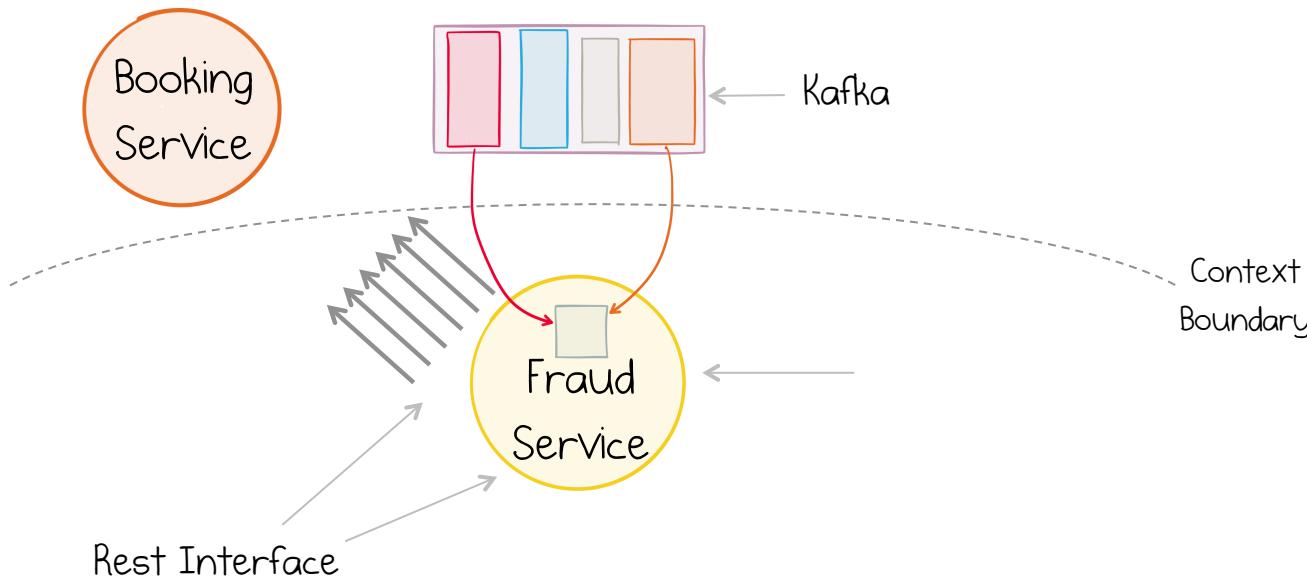


Pattern 3: External DB

- Again similar, but heavier weight
- Use when you cannot pre-compute the view (i.e. to do ad hoc queries)
- Trick: start with only the data (fields) you need today. You can always go back for more.



Pattern 4: Hybrid



NB this breaks the async model, but can
be a useful compromise in some cases



Pattern 4: Hybrid

- Quite common in practice
- Allows you to break out of the async world
- Use sparingly, but do use when appropriate (e.g. login service)



Benefits

- Forms a central, immutable narrative
- Communication Protocol and Event Store become one
- Better decoupling (RDFC)
- Excellent scalability
- Several approaches for managing queries, lightweight => heavy weight
- Event sourcing & CQRS at its core



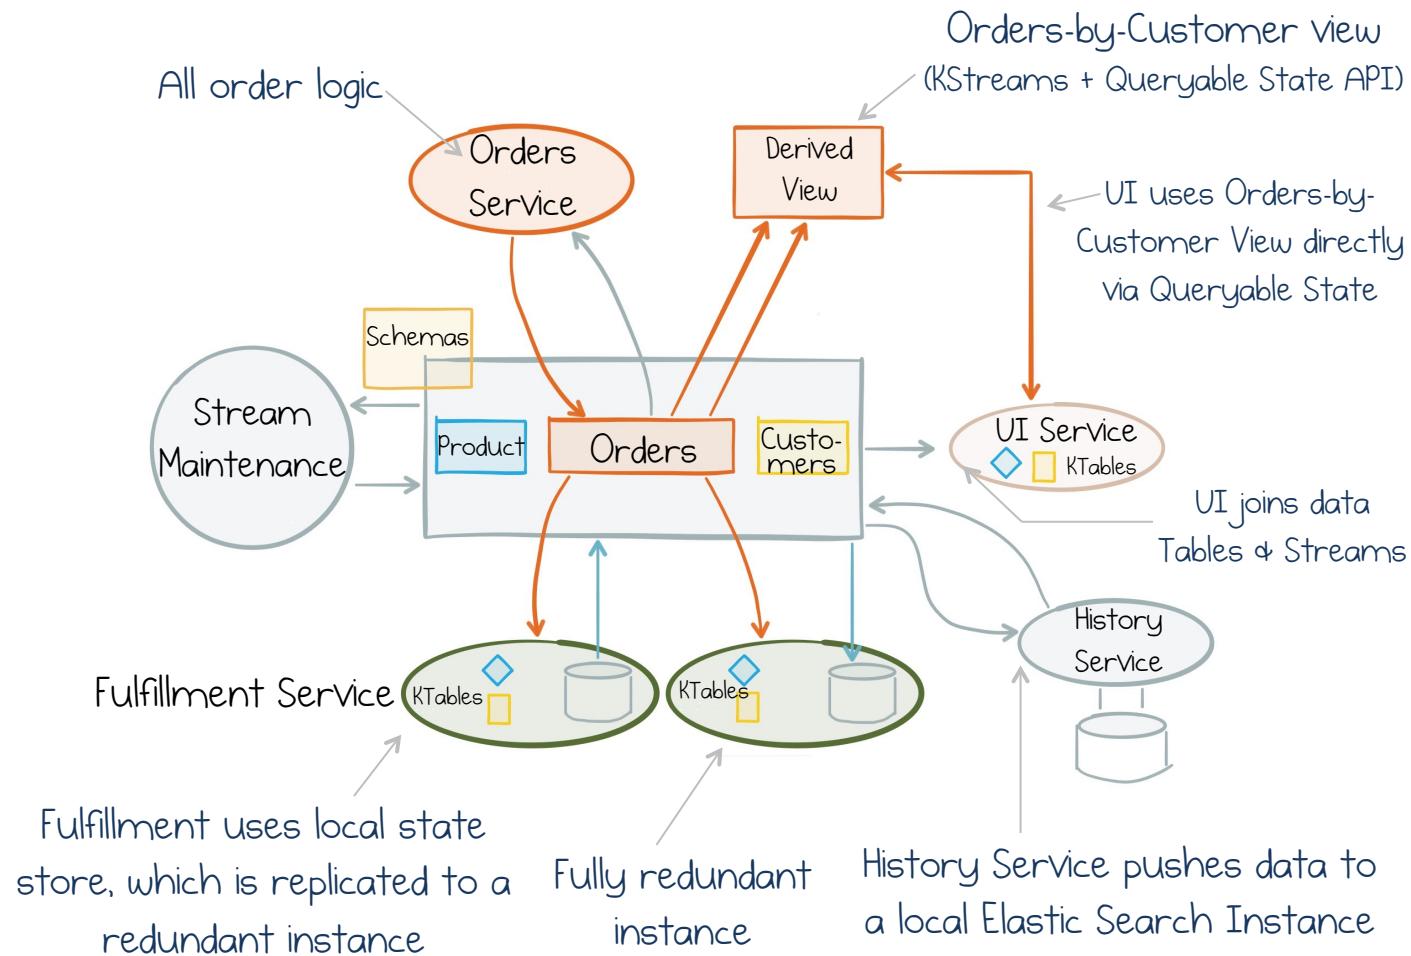
INTERACT THROUGH
EVENTS.
DON'T TALK TO
SERVICES



BUT IT'S OK TO BREAK
THE RULES!



Pulling it all together





Stay in touch!

Online Talks

cnfl.io/online-talks

